

PAS_TIX User's manual

November 22, 2012

Contents

1	About PaStiX	4
1.1	Introduction	4
1.2	Ordering	4
1.3	Symbolic factorization	6
1.4	Distribution and scheduling	6
1.5	Factorization	7
1.6	Solve	7
1.7	Refinement	7
1.7.1	GMRES	7
1.7.2	Conjuguate gradiant	7
1.7.3	Simple iterative refinement	7
1.8	Out-of-core	8
2	Compilation of PaStiX	9
2.1	Quick compilation of PASTIX	9
2.1.1	Pre-requirement	9
2.1.2	compilation	9
2.2	Makefile keywords	10
2.3	Compilation flags	10
2.3.1	Integer types	10
2.3.2	Coefficients	10
2.3.3	MPI and Threads	10
2.3.4	Ordering	11
2.3.5	NUMA aware library	11
2.3.6	Statistics	11
2.3.7	Out-of-core	11
2.3.8	Python wrapper	12
3	PaStiX options	13
3.1	Global parameters	13
3.1.1	Verbosity level	13
3.1.2	Indicating the steps to execute	13
3.1.3	Symmetry	13
3.1.4	Threads	14
3.1.5	Matrix descriptrion	14
3.2	Initialisation	14
3.2.1	Description	14
3.2.2	Parameters	15
3.3	Ordering	15
3.3.1	Description	15
3.3.2	Parameters	15

3.4	Symbolic factorisation	16
3.4.1	Description	16
3.4.2	Parameters	16
3.5	Analyse	16
3.5.1	Description	16
3.5.2	Parameters	16
3.5.3	Output	16
3.6	Numerical factorisation	16
3.6.1	Description	16
3.6.2	Parameters	17
3.6.3	Ouputs	17
3.7	Solve	17
3.7.1	Description	17
3.7.2	Parameters	17
3.7.3	Ouputs	17
3.8	Refinement	17
3.8.1	Description	17
3.8.2	Parameters	18
3.8.3	Output	18
3.9	Clean	18
4	PaStiX functions	19
4.1	Original sequential interface	19
4.1.1	The main function	19
4.2	Original Distributed Matrix Interface	20
4.2.1	A distributed CSC	20
4.2.2	Usage of the distributed interface	21
4.2.3	The distributed interface prototype	21
4.3	Murge : Uniformized Distributed Matrix Interface	21
4.3.1	Description	21
4.3.2	Additional specific functions for PASTiX	23
4.4	Auxiliary PASTiX functions	27
4.4.1	Distributed mode dedicated functions	27
4.4.2	Binding thread	28
4.4.3	Working on CSC or CSCD	29
4.4.4	Schur complement	34
4.5	Multi-arithmetic	37
5	SWIG python wrapper	39
5.1	Requirement	39
5.2	Building the python wrapper	39
5.3	Example	39
6	Examples	40
6.1	Examples in C	40
6.1.1	Parameters	40
6.1.2	simple.c	41
6.1.3	simple_dist.c	41
6.1.4	do_not_redispatch_rhs.c	41
6.1.5	step-by-step.c	41
6.1.6	step-by-step_dist.c	41
6.1.7	reentrant.c	41
6.1.8	multi-comm.c	42
6.1.9	multi-comm-step.c	42

6.1.10	<code>schur.c</code>	42
6.1.11	<code>isolate_zeros.c</code>	42
6.2	Examples in Fortran	42
6.2.1	Parameters	42
6.2.2	<code>fsimple.f90</code>	42
6.2.3	<code>fstep-by-step.f90</code>	42
6.2.4	<code>Murge-Fortran.f90</code>	43
7	Outputs	44
7.1	Logs	44
7.1.1	Controlling output logs	44
7.1.2	Interpreting output logs	44
7.2	Integer and double outputs	46
7.2.1	Integer parameters	46
7.2.2	Double parameters	46

Chapter 1

About PaStiX

1.1 Introduction

PASTIX is a complete, parallelized, and multi-threaded library for the resolution of huge linear systems of equations. It is developed by BACCHUS team from INRIA¹.

Depending on characteristics of the matrix A , the solution of $Ax = b$ can be computed in several ways :

- if the matrix A is symmetric positive-definite, we can use the Cholesky ($A = LL^t$, L lower triangular matrix, L^t its transpose) or Cholesky-Crout ($A = LDL^t$, D diagonal matrix) with, or without, numerical pivoting,
- if the matrix A is not symmetric, the LU decomposition (U upper triangular matrix) with static pivoting will be used.

PASTIX steps :

- Reordering the unknowns in order to reduce the fill-in induced by the decomposition,
- Symbolic factorization, to predict the structure of the factorized matrix,
- Distributing matrix blocks among the processors,
- Decomposition of the matrix A ,
- Solving the system (up-down),
- Refining the solution because we use static pivoting, which can reduce the precision of the solution.

1.2 Ordering

PASTIX computes the reordering by calling the SCOTCH package. (METIS can also be used.) During direct decomposition, new nonzero terms, called “fill-in”, appear in the decomposed matrix.

In a graph $G(V, E)$, whose vertices are the unknowns, and whose edges are defined by : $(i, j) \in E(G) \leftrightarrow a_{ij} \neq 0$.

¹Institut National de Recherche en informatique et Automatique.

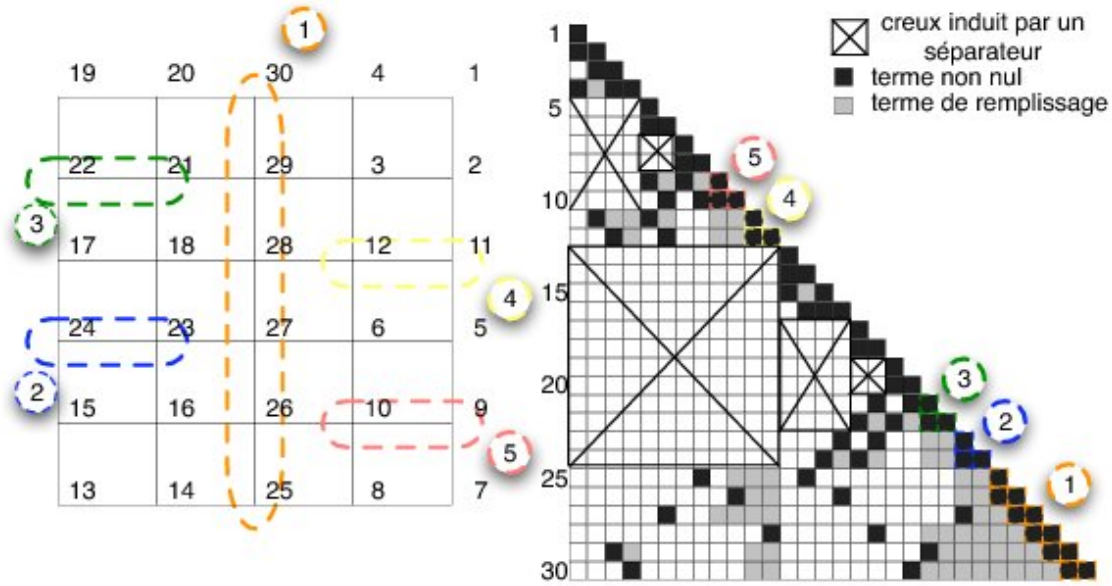


Figure 1.1: Nested dissection : Left, the graph before fill-in; right, the matrix with fill-in.

One fill-in edge will be created between two unknowns i and j if there is a path from i to j passing only through vertices with number lower than $\min(i, j)$.

To save memory and computation time during decomposition, this fill-in has to be minimized. To manage this several algorithms are used.

First, we use a nested dissection (fig. 1.1, p. 5) : we search in the graph for a separator S of minimal size that cuts the graph into two parts of about the same size.

Then the separator nodes are indexed with the highest numbers available, so that no fill-in terms will appear during decomposition between the unknowns of the two separated parts of the graph. The algorithm is then repeated on each subgraph. This step also produces a dependency graph used to distribute computations onto the processors.

When the subgraph is smaller than a specified threshold, the Halo Approximate Minimum Degree algorithm (fig. 1.2, p. 6) is used.

This algorithm further reduces the number of “fill-in” terms even more by assigning the smallest available numbers to the nodes with lowest degree. Direct and foreign neighbors (halo) are considered, in order not to neglect interactions with the whole graph.

An elimination tree can be constructed by using the following rule : there is an edge between i and j if the first nonzero entry of column i is at row j .

This tree represents the dependencies between the computation steps. The wider and deeper it is, the less the computation are dependent from each other.

The goal of the reordering step is to minimize “fill-in” and to reduce the dependencies between computations.

This algorithm also compute a partition of the columns. This partition results from the fusion of the separators and the sub-graphs reordered using the Halo Approximate Minimum Degree algorithm.

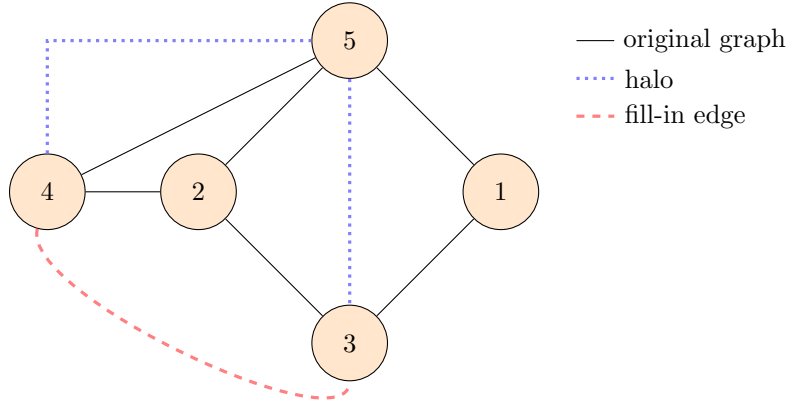


Figure 1.2: Halo Approximate Minimum Degree.

1.3 Symbolic factorization

The symbolic factorization step computes the structure of the factorized matrix from the original matrix A .

The factorization is done by blocks. This computation is cheap, its complexity depend on the number of extra-diagonal blocs in the decomposed matrix.

1.4 Distribution and scheduling.

PASTIX uses static ordering by default. Computation is organized, in advance, for maximum alignment with the parallel architecture being used.

The algorithm is composed of two steps.

The first step is the partitioning step. The largest blocks are cut into smaller pieces to be computed by several processors. *Full block parallel computation* parallelism is used. It is applied on biggest separators of the elimination tree, after the decomposition.

During this step, the elimination tree is traversed from the root to the leaves and for each block, a set of processor candidates able to handle it is computed. This idea of “processor candidate” is essential for the preservation of communication locality; this locality is obtained following the elimination tree, because only the processors which compute a part of the subtree of a node would be candidate for this node.

The second step is the distributing phase, each block will be associated to one of its candidates which will compute it.

To do so, the elimination tree is climbed up from the leaves to the root and, to distribute every node of the tree on processors, the computation and communication time is computed for each block. Thus, each block is allocated to the processor which will finish its processing first.

This step schedules computation and communication precisely; it needs an accurate calibration of BLAS3 and communications operations on the destination architecture. The perf module contains information about this calibration.

Thus, three levels of parallelism are used :

- Coarse grained parallelism, induced by independent computation between two subtrees of one node. This parallelism is also called sparse induced parallelism,

- Medium grained parallelism, induced by dense blocks decomposition. This is a node level parallelism, due to the possibility of cutting elimination tree nodes to distribute them onto different processors,
- thin grained parallelism, obtained by using BLAS3 operations. This requires that the block size be correctly chosen.

1.5 Factorization

The numeric solution of the problem computes a parallel Crout (LL^t or LDL^t) or Cholesky (LU) decomposition with a one-dimension column block distribution of the matrix.

Two type of algorithm exist, *super-nodal*, coming directly from column elimination algorithm rewritten by blocks, and *multi-frontal*.

Among *super-nodals* methods, the *fan-in* or *left-looking* method involves getting previous column blocks modifications to compute next column blocks.

The contrary, the *fan-out* or *right-looking* method involves computing the current column block and modifying next column blocks.

In PASTIX, the *fan-in* method, more efficient in parallel, has been implemented. This method has the advantage of considerably reducing communication volume.

1.6 Solve

The solve step uses the up-down method. Since this computation is really cheap compared to decomposition the decomposition distribution is kept.

1.7 Refinement

If the precision of the result of the Solve step is insufficient, iterative refinement can be performed. This section described the several iterative methods implemented in PASTIX.

1.7.1 GMRES

The GMRES algorithm used in PASTIX is based on work with Youssef Saad that can be found in [Saa96].

1.7.2 Conjugate gradient

The conjugate gradient method is available with LL^t and LDL^t factorizations.

1.7.3 Simple iterative refinement

This method, described in [algorithm 1 p.8](#), computes the difference b_i between b and Ax and solves the system $Ax = b_i$. The solution x_i are added to the initial solution, and the process repeated until the relative error ($Max_k(\frac{|b-Ax|_k}{||A||x|+|b||_k})$) is small enough.

This method is only available with LU factorization.

Algorithm 1 Static pivoting refinement algorithm

```
lerr  $\leftarrow$  0;  
flag  $\leftarrow$  TRUE;  
while flag do  
  r  $\leftarrow$  b − A × x;  
  r'  $\leftarrow$  |A||x| + |b|;  
  err  $\leftarrow$   $\max_{i=0..n}(\frac{r[i]}{r'[i]});$   
  if last_err = 0 then  
    | last_err  $\leftarrow$  3 × err;  
  rberror  $\leftarrow$  ||r||/||b||;  
  if (iter < iparm[IPARM_ITERMAX] and  
    err > dparm[DPARM_EPSILON_REFINEMENT] and  
    err  $\leq$   $\frac{lerr}{2}$ )  
  then  
    | r'  $\leftarrow$  x;  
    | x  $\leftarrow$  solve(A × x = r);  
    | r'  $\leftarrow$  r' + x;  
    | last_err  $\leftarrow$  err;  
    | iter  $\leftarrow$  iter + 1;  
  else  
    | flag  $\leftarrow$  FALSE;
```

1.8 Out-of-core

An experimental out-of-core version of PASTIX has been written.

Directions for compiling this version are given in section 2.3.7.

For the moment, only multi-threaded version is supported, the hybrid “MPI+threads” version is likely to lead to unresolved deadlocks.

The disk inputs and outputs are managed by a dedicated thread. This thread will prefetch column blocks and communication buffers so that computing threads will be able to use them.

As the whole computation as been predicted, the out-of-core thread will follow the scheduled computation and prefetch needed column blocks. It will also save them, depending on their next acces, if the memory limit is reached.

The prefetch algorithm is not dependant of the number of computation threads, it will follow the computation algorithm as if there was only one thread.

Chapter 2

Compilation of PaStiX

This chapter will present how to quickly compile PASTiX and the different compilation options that can be used with PASTiX.

2.1 Quick compilation of PaStiX

2.1.1 Pre-requirement

To compile PASTiX, a BLAS library is required.

To compile PASTiX, it is advise to get first SCOTCH or PT-SCOTCH ordering library (<https://gforge.inria.fr/projects/scotch/>).

However, it is possible to compile PASTiX with **Metis** or without any ordering (using user ordering), or even both. **Metis** and SCOTCH or PT-SCOTCH.

To be able to use threads in PASTiX, the **pthread** library is required.

For a MPI version of PASTiX, a MPI library is obviously needed.

2.1.2 compilation

To compile PASTiX, select in **src/config/** the compilation file corresponding to your architecture, and copy it to **src/config.in**.

You can edit this file to select good libraries and compilation options.

Then you can run **make expor install** to compile PASTiX library.

This compilation will produce PASTiX library, named **libpastix.a**; PASTiX C header, named **pastix.h**; a Fortran header named **pastix_fortran.h** (use it with **#include**) and a script, **pastix-conf** that will describes how PASTiX has been compiled (options, flags...). This script is used to build the Makefile in **example/src**.

Another library is produced to use **Murge** interface : **libpastix_murge.a**, which works with the C header **murge.h** and the Fortran header **murge.inc** (a true Fortran include).

2.2 Makefile keywords

make help : print this help;
make all : build PASTIX library;
make debug : build PASTIX library in debug mode;
make drivers : build matrix drivers library;
make debug drivers : build matrix drivers library in debug mode;
make examples : build examples (will run 'make all' and 'make drivers' if needed);
make murge : build murge examples (only available in distributed mode `-DDISTRIBUTED`, will run 'make all' and 'make drivers' if needed);
make python : Build python wrapper and run `src/simple/pastix-python.py`;
make clean : remove all binaries and objects directories;
make cleanall : remove all binaries, objects and dependencies directories

2.3 Compilation flags

2.3.1 Integer types

PASTIX can be used with different integer types. User can choose the integer type by setting compilation flags.

The flag `-DINTSIZE32` will set PASTIX integers to 32 bits integers.

The flag `-DINTSIZE64` will set PASTIX integers to 64 bits integers.

If you are using Murge interface, you can also set `-DINTSSIZE64` to set Murge's INTS integers to 64 bits integers. This is not advised, INTS should be 32 bit integers.

2.3.2 Coefficients

Coefficients can be defined to several types:

real : using no flag,

double : using flag `-DPREC_DOUBLE`,

complex : using flag `-DTYPE_COMPLEX`,

double complex : using both `-DTYPE_COMPLEX` and `-DPREC_DOUBLE` flags.

2.3.3 MPI and Threads

PASTIX default version uses threads and MPI.

The expected way of running PASTIX is with one MPI process by node of the machine and one thread for each core.

It is also possible to deactivate MPI using `-DFORCE_NOMPI` and threads using `-DFORCE_NOSMP`. User only has to uncomment the corresponding lines of his `config.in` file.

PASTiX also proposes the possibility to use one thread to handle communications reception. It can give better results if the MPI library does not handle correctly the communication progress. This option is activated using `-DTHREAD_COMM`.

If the MPI implementation does not handle `MPI_THREAD_MULTIPLE` a funneled version of PASTiX is also proposed. This version, available through `-DTHREAD_MULTIPLE`, can affect the solver performances.

An other option of PASTiX is provided to suppress usage of MPI types in PASTiX if the MPI implementation doesn't handle it well. This option is available with the compilation flag `-DNO_MPI_TYPE`.

The default thread library used by PASTiX is the `POSIX` one.

The `MARCEL` library, from the `INRIA` team `RUNTIME` can also be used. Through `marcel`, the `Bubble` scheduling framework can also be used defining the option `-DPASTIX_USE_BUBBLE`.

All this options can be set in the `MARCEL` section of the `config.in` file.

2.3.4 Ordering

The graph partitioning can be done in PASTiX using `SCOTCH`, `PT-SCOTCH` or `METIS`. It can also be computed by user and given to PASTiX through the permutation arrays parameters. To activate the possibility of using `SCOTCH` in PASTiX (default) uncomment the corresponding lines of the `config.in`.

In the same way, to use `METIS`, uncomment the corresponding lines of the `config.in` file.

`SCOTCH` and `METIS` can be used together, alone, or can be unused.

`PT-SCOTCH` is required when the `-DDISTRIBUTED` flag has been set, that is to say, when compiling with the distributed interface.

When using centralised interface with `PT-SCOTCH`, the ordering will be performed with `SCOTCH` functions.

2.3.5 NUMA aware library

To be more efficient on NUMA machines, the allocation of the matrix coefficient is can be performed on each thread.

The default compilation flag `-DNUMA_ALLOC` will activate this “per thread” allocation.

2.3.6 Statistics

Memory usage

The compilation flag `MEMORY_USAGE` can be used to display memory usage at deferent steps of PASTiX run.

The information also appear in double parameter output array, at index `DPARM_MEM_MAX`. The value is given in octets.

2.3.7 Out-of-core

It is possible to experiment an out-of-core version of PaStiX.

To compile PASTiX with Out Of Core, compile it with `-DOOC`.

2.3.8 Python wrapper

A simple python wrapper can be built to use PASTIX from python.

This wrapper uses SWIG (Simplified Wrapper and Interface Generator - www.swig.org).

To build the python interface, user needs to use dynamic (or at least built with -fPIC) libraries (for BLAS, MPI and SCOTCH).

It is also necessary to uncomment the `Python` part of the `config.in` file and set paths correctly for `MPI4PY_DIR`, `MPI4PY_INC` and `MPI4PY_LIBDIR`.

Then you just have to run `make python` to build the interface and test it with `examples/src/pastix_python.py`.

Chapter 3

PaStiX options

PASTiX can be called step by step or in one unique call.
User can refer to PASTiX step-by-step and simple examples.
The following section will describe each steps and options that can be used to tune the computation.

3.1 Global parameters

This section present list of parameters used in several PASTiX steps.

3.1.1 Verbosity level

IPARM_VERBOSE : used to set verbosity level. Can be set to 3 values :

- API_VERBOSE_NOT : No display at all,
- API_VERBOSE_NO : Few displays,
- API_VERBOSE_YES : Maximum verbosity level.

3.1.2 Indicating the steps to execute

IPARM_START_TASK : indicates the first step to execute (cf. quick reference card).
Should be set before each call.
It is modified during PASTiX calls, at the end of one call it is equal to the last step performed
(should be IPARM_END_TASK).

IPARM_END_TASK : indicates the last step to execute (cf. quick reference card).
Should be set before each call.

NB : Setting IPARM_MODIFY_PARAMETER to API_NO will make PASTiX initialize integer and floating point parameters. After that operation, PASTiX will automatically return, without taking into account IPARM_START_TASK nor IPARM_END_TASK.

3.1.3 Symmetry

IPARM_FACTORISATION : Gives the factorisation type.
It can be API_FACT_LU, API_FACT_LDLT or API_FACT_LLT.

It has to be set from ordering step to refinement step to the same value.

With non symmetric matrix, only *LU* factorisation is possible.

IPARM_SYM : Indicates if the matrix is symmetric or not.

With symmetric matrix, only the inferior triangular part has to be given.

3.1.4 Threads

To save memory inside a SMP node users can use threads inside each node.

Each thread will allocate the part of the matrix he will work mostly on but all threads will share the matrix local to the SMP node.

IPARM_THREAD_NBR : Number of computing threads per MPI process,

IPARM_BINDTHRD : can be set to the mode used to bind threads on processors :

API_BIND_NO : do not bind threads on processors,

API_BIND_AUTO : default binding mode (thread are binded cyclicaly (thread n to proc $\lfloor \frac{n}{procnbr} \rfloor$),

API_BIND_TAB : Use vector given by `pastix_setBind (?? p.??)`.

This section describes which steps are affected by which options.

3.1.5 Matrix description

PASTIX solver can handle different type of matrices.

User can describe the matrix using several parameters :

IPARM_DOF_NBR : indicate the number of degree of freedom by edge of the graph. The default value is one.

IPARM_SYM : indicate if the matrix is symmetric or not. This parameters can be set to two values :

API_SYM_YES : If the matrix is symmetric.

API_SYM_NO : If the matrix is not symmetric.

If user is not sure that the matrix will fit PASTIX and Scotch requirements, the parameters **IPARM_MATRIX_VERIFICATION** can be set to **API_YES**.

With distributed interface, to prevent PASTIX from checking that the matrix has been correctly distributed after distribution computation, **IPARM_CSCD_CORRECT** can be set to **API_YES**.

3.2 Initialisation

3.2.1 Description

This steps initializes the `pastix_data` structure for next PASTIX calls.

It can also initialize integer and double parameters values (see quick reference card for default values).

It has to be called first.

3.2.2 Parameters

To call this step you have to set `IPARM_START_TASK` to `API_TASK_INIT`.

When this step is called, `IPARM_MODIFY_PARAMETER` should be set to `API_NO`. This will make PASTIX set all integer and double parameters.

If `IPARM_MODIFY_PARAMETER` is set to `API_NO`, `pastix` will automatically return after initialisation, whatever `IPARM_END_TASK` is set to.

The user CSC matrix can be checked during this step. To perform matrix verification, `IPARM_MATRIX_VERIFICATION` has to be set to `API_YES`.

This will correct numbering if the CSC is in C numbering, sort the CSC and check if the graph of the matrix is symmetric (only if the matrix is non-symmetric (depending on `iparm[IPARM_SYM]` value)).

It is also possible to use `pastix_checkMatrix` function ([4.4.3 p.29](#)) to perform this checking operations and also correct the graph symmetry .

3.3 Ordering

3.3.1 Description

The ordering step will reorder the unknowns of the matrix to minimize fill-in during factorisation. This step is description in [1.2](#)

3.3.2 Parameters

To call ordering step, `IPARM_START_TASK` has to be set to `API_TASK_ORDERING` or previous task and `IPARM_END_TASK` must be greater or equal to ordering one.

Ordering can be computed with `SCOTCH` or `Metis`.

To enable `Metis`, user has to uncomment the corresponding part of the `config.in` file. To select the ordering software, user may set `IPARM_ORDERING` to :

- `API_ORDER_SCOTCH` : use `Scotch` for the ordering (default)
- `API_ORDER_METIS` : use `Metis` for the ordering.

To have a finer control on the ordering software, user can set `IPARM_DEFAULT_ORDERING` to `API_YES` and modify those parameters:

- when using `Scotch` :
 - `IPARM_ORDERING_SWITCH_LEVEL` : ordering switch level (see `Scotch User's Guide`),
 - `IPARM_ORDERING_CMIN` : ordering cmin parameter (see `Scotch User's Guide`),
 - `IPARM_ORDERING_CMAX` : ordering cmax parameter (see `Scotch User's Guide`),
 - `IPARM_ORDERING_FRAT` : ordering frat parameter ($\times 100$) (see `Scotch User's Guide`).
- when using `Metis` :
 - `IPARM_ORDERING_SWITCH_LEVEL` : `Metis ctype` option,
 - `IPARM_ORDERING_CMIN` : `Metis itype` option,
 - `IPARM_ORDERING_CMAX` : `Metis rtype` option,
 - `IPARM_ORDERING_FRAT` : `Metis dbgvl` option,

IPARM_STATIC_PIVOTING : Metis oflags option,
IPARM_METIS_PFACTOR : Metis pfactor option,
IPARM_NNZERO : Metis nseps option.

3.4 Symbolic factorisation

3.4.1 Description

This step is the symbolic factorisation step described in [1.3](#).

3.4.2 Parameters

If user didn't called ordering step with *Scotch*, IPARM_LEVEL_OF_FILL has to be set to -1 to use KASS algorithm.

If PASTiX was compiled with -DMETIS, this parameter will be forced to -1.

If PASTiX runs KASS algorithm, IPARM_AMALGAMATION_LEVEL will be take into account. KASS will merge supernodes untill fill reaches IPARM_AMALGAMATION_LEVEL.

3.5 Analyse

3.5.1 Description

During this step, PASTiX will simulate factorization and decide where to assign each part of the matrix.

More details can be read in [1.4](#)

3.5.2 Parameters

IPARM_MIN_BLOCKSIZE : Minimum size of the column blocks computed in blend;

IPARM_MAX_BLOCKSIZE : Maximum size of the column blocks computed in blend.

3.5.3 Output

DPARM_ANALYZE_TIME : time to compute analyze step,

DPARM_PRED_FACT_TIME : predicted factorization time (with IBM PWR5 ESSL).

3.6 Numerical factorisation

3.6.1 Description

Numerical factorisation ([1.5](#)) of the given matrix.

Can be LU , LL^t , or LDL^t factorisation.

3.6.2 Parameters

DPARM_EPSILON_MAGN_CTRL : value which will be used for static pivoting. Diagonal values smaller than it will be replaced by it.

3.6.3 Ouputs

After factorisation, **IPARM_STATIC_PIVOTING** : will be set to the number of static pivoting performed.

A static pivoting is performed when a diagonal value is smaller than **DPARM_EPSILON_MAGN_CTRL**.

IPARM_ALLOCATED_TERMS : Number of non zeros allocated in the final matrix,

DPARM_FACT_TIME : contains the time spent computing factorisation step in seconds, in real time, not cpu time,

DPARM_SOLV_FLOPS : contains the number of operation per second during factorisation step.

3.7 Solve

3.7.1 Description

This step, described in 1.6, will compute the solution of the system.

3.7.2 Parameters

IPARM_RHS_MAKING : way of obtaining the right-hand-side member :

API_RHS_B : get right-hand-side member from user,

API_RHS_1 : construct right-hand-side member such as the solution X is defined by : $\forall i, X_i = 1$,

API_RHS_I : construct right-hand-side member such as the solution X is defined by : $\forall i, X_i = i$,

3.7.3 Ouputs

DPARM_SOLV_TIME : contains the time spent computing solve step, in second, in real time, not cpu time,

DPARM_SOLV_FLOPS : contains the number of operation per second during solving step.

3.8 Refinement

3.8.1 Description

After solving step, it is possible to call for refinement if the precision of the solution is not sufficient.

3.8.2 Parameters

To call for refinement, `IPARM_START_TASK` must be lower than `API_TASK_REFINEMENT` and `IPARM_END_TASK` must be greater than `API_TASK_REFINEMENT`.

A list of parameters can be used to setup refinement step :

`IPARM_ITERMAX` : Maximum number of iteration in refinement step,

`IPARM_REFINEMENT` : Type of refinement :

`API_RAF_GMRES` : GMRES algorithm,

`API_RAF_PIVOT` : a simple iterative algorithm, can only be used with LL^t or LDL^t factorization (the corresponding `iparm` must be correctly set),

`API_RAF_GRAD` : conjugate gradient algorithm, can only be used with LU factorization (the corresponding `iparm` must be correctly set).

`IPARM_GMRES_IM` : syze of the Krylov space used in GMRES,

`DPARAM_EPSILON_REFINEMENT` : Desired solution precision.

3.8.3 Output

After refinement, `IPARM_NBITER` will be set to the number of iteration performed during refinement.

The value `DPARAM_RELATIVE_ERROR` will be the error between AX and B . It should be smaller than `DPARAM_EPSILON_REFINEMENT`.

`DPARAM_RAFF_TIME` contains the time spent computing solve step, in second.

3.9 Clean

This step simply free all memory used by PaStiX.

Chapter 4

PaStiX functions

The PASTiX library provides several functions to setup and run PASTiX decomposition steps. Two different ways of using PASTiX exist, it can be called with a sequential matrix in input, or with a distributed matrix.

The sequential library is composed of a main function **pastix**. The distributed PASTiX library has two interfaces. The original one is based on the sequential one, it is composed of few auxiliary functions and one main function which, depending on parameters, will run all steps.

The third one is an interface which has been adapted to match with HIPS and MUMPS interfaces.

4.1 Original sequential interface

The original interface is composed of one main function. Few auxiliary functions also permit to check that the user matrix will fit with PASTiX matrix format.

4.1.1 The main function

The main function of the original sequential interface is the **pastix** function (Fig. 4.1, p.19).

It is used to run separately or in one call all PASTiX decomposition steps.

A Fortran interface to this function is also available in PASTiX library (Fig. 4.2, p.20).

In this centralised interface to PASTiX, all parameters should be equal on all MPI processors.

The first parameter, **pastix_data** is the address of a structure used to store data between **pastix** calls. It has to be set to NULL (0 in Fortran) before first call. This parameter is modified by PASTiX and should be untouched until the end of the program execution.

```
#include "pastix.h"
void pastix ( pastix_data_t ** pastix_data, MPI_Comm      pastix_comm,
              pastix_int_t   n,          pastix_int_t * colptr,
              pastix_int_t * row,         pastix_float_t * avals,
              pastix_int_t * perm,        pastix_int_t * invp,
              pastix_float_t * b,         pastix_int_t  rhs,
              pastix_int_t * iparm,       double        * dparm );
```

Figure 4.1: PASTiX main function prototype

```

#include "pastix_fortran.h"
pastix_data_ptr_t :: pastix_data
integer           :: pastix_comm
pastix_int_t      :: n, rhs, ia(n), ja(nnz)
pastix_float_t    :: avals(nnz), b(n)
pastix_int_t      :: perm(n), invp(n), iparm(64)
real*8            :: dparm(64)
call pastix_fortran ( pastix_data, pastix_comm, n, ia, ja, avals,
                      perm, invp, b, rhs, iparm, dparm )

```

Figure 4.2: PASTIX main function fortran interface

$$\begin{pmatrix} 1 & 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 & 0 \\ 2 & 0 & 5 & 0 & 0 \\ 0 & 4 & 6 & 7 & 0 \\ 0 & 0 & 0 & 0 & 8 \end{pmatrix}
\begin{array}{lcl} \text{colptr} & = & \{1, 3, 5, 7, 8, 9\} \\ \text{row} & = & \{1, 3, 2, 4, 3, 4, 4, 5\} \\ \text{avals} & = & \{1, 2, 3, 4, 5, 6, 7, 8\} \end{array}$$

Figure 4.3: CSC matrix example

`pastix_comm` is the MPI communicator used in PASTIX.

`n`, `colptr`, `row` and `avals` is the matrix to factorize, in CSC representation (Fig. 4.3 p.20).

`n` is the size of the matrix, `colptr` is an array of $n + 1$ `pastix_int_t`. It contains index of first elements of each column in `row`, the row of each non null element, and `avals`, the value of each non null element.

`perm` (resp. `invp`) is the permutation (resp. reverse permutation) tabular. It is an array of `n` `pastix_int_t` and must be allocated by user. It is set during ordering step but can also be set by user.

`b` is an array of $n \times \text{rhs}$ `pastix_float_t`. It correspond to the right-hand-side member(s) and will contain the solution(s) at the end of computation.

Right-hand-side members are contiguous in this array.

`rhs` is the number of right-hand-side members. Only one is currently accepted by PASTIX.

`iparm` is the integer parameters array, of `IPARM_SIZE` `pastix_int_t`.

`dparm` is the floating parameters array, of `DPARM_SIZE` double.

The only parameters that can be modified by this function are `pastix_data`, `iparm` and `dparm`.

4.2 Original Distributed Matrix Interface

A new interface was added to run PASTIX using distributed data (Fig. 4.6 p.22 and Fig. 4.7 p.22).

The Data is given distributed by columns.

The original CSC is replaced by a distributed CSC (Fig. 4.4, p.21).

4.2.1 A distributed CSC

A distributed CSC is a CSC with a given list of columns. An additionnal array give the global number of each local column.

dCSC matrix example :

$$\begin{pmatrix} P_1 & P_2 & P_1 & P_2 & P_1 \\ 1 & 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 & 0 \\ 2 & 0 & 5 & 0 & 0 \\ 0 & 4 & 6 & 7 & 0 \\ 0 & 0 & 0 & 0 & 8 \end{pmatrix}$$

On processor one :

colptr = {1, 3, 5, 6}
row = {1, 3, 3, 4, 5}
avals = {1, 2, 5, 6, 8}
loc2glb = {1, 3, 5}

On processor two :

colptr = {1, 3, 4}
row = {2, 4, 4}
avals = {3, 4, 7}
loc2glb = {2, 4}

Figure 4.4: A distributed CSC

4.2.2 Usage of the distributed interface

A good usage of the distributed interface would follow this steps (Fig. 4.5 p.22) :

1. provide the graph to PASTIX in user's distribution to perform analyze steps;
2. get the solver distribution from PASTIX.
3. provide the matrice and right-hand-side in PASTIX distribution to avoid redistribution inside the solver and perform factorization and solving steps.

An example of this utilisation of the distributed interface can be found in `src/example/src/simple_dist.c`.

4.2.3 The distributed interface prototype

The distributed interface prototype is similar to the centralised one, only a `loc2glob` array is added to it.

4.3 Murge : Uniformized Distributed Matrix Interface

4.3.1 Description

This interface was added to PASTIX to simplify the utilisation of multiple solvers, HIPS and PASTIX in a first step, other Murge compliant solvers later.

It is composed of a large number of function but it is less flexible than the original one.

Using this interface, you can change between solvers with only few modifications. You just have to set specific solvers option.

This new interface has been created trying to simplify user work, thinking about his needs.

The graph is simply built in order to compute the distribution of the problem. Then the matrix is filled taking account or ignoring the solver distribution. After that the right-hand-side id given and the solution is computed.

More information about this new interface can be found at <http://murge.gforge.inria.fr/>.

```

/* Build the CSCd graph with external software distribution */

iparm[IPARMSTART_TASK]          = APLTASK_ORDERING;
iparm[IPARMEND_TASK]            = APLTASK_BLEND;

dpastix(&pastix_data , MPLCOMM_WORLD,
        ncol, colptr, rows, NULL, loc2glob,
        perm, NULL, NULL, 1, iparm, dparm);

ncol2 = pastix_getLocalNodeNbr(&pastix_data);
if (NULL == (loc2glob2 = malloc(ncol2 * sizeof(pastix_int_t))))
{
    fprintf(stderr, "Malloc error\n");
    return EXIT_FAILURE;
}

pastix_getLocalNodeLst(&pastix_data, loc2glob2);

... /* Building the matrix following PaStiX distribution */

iparm[IPARMSTART_TASK]          = APLTASK_NUMFACT;
iparm[IPARMEND_TASK]            = APLTASK_CLEAN;

dpastix(&pastix_data , MPLCOMM_WORLD,
        ncol2, colptr2, rows2, values2, loc2glob2,
        perm, invp, rhs, 1, iparm, dparm);

```

Figure 4.5: Using distributed interface

```

#include "pastix.h"
void dpastix ( pastix_data_t ** pastix_data, MPI_Comm  pastix_comm,
               pastix_int_t  n,          pastix_int_t * colptr,
               pastix_int_t  * row,      pastix_float_t * avals,
               pastix_int_t  * loc2glb,
               pastix_int_t  * perm,     pastix_int_t * invp,
               pastix_float_t * b,       pastix_int_t  * rhs,
               pastix_int_t  * iparm,    double      * dparm );

```

Figure 4.6: Distributed C interface

```

#include "pastix_fortran.h"
pastix_data_ptr_t  :: pastix_data
integer            :: pastix_comm
pastix_int_t       :: n, rhs, ia(n+1), ja(nnz)
pastix_float_t     :: avals(nnz), b(n)
pastix_int_t       :: loc2glb(n), perm(n), invp(n), iparm(64)
real*8             :: dparm(64)
call dpastix_fortran ( pastix_data, pastix_comm, n, ia, ja, avals,
                      loc2glob perm, invp, b, rhs, iparm, dparm )

```

Figure 4.7: Distributed fortran interface

4.3.2 Additional specific functions for PaStiX

Few auxiliary functions were added in the PASTiX implementation of this interface. They are not essential, Murge can be used without this functions.

MURGE_Analyze

```
INTS MURGE_Analyze ( INTS id );
```

```
SUBROUTINE MURGE_ANALYZE ( ID, IERROR)
  INTS, INTENT(IN)      :: ID
  INTS, INTENT(OUT)     :: IERROR
END SUBROUTINE MURGE_ANALYZE
```

Parameters :

id : Solver instance identification number.

Perform matrix analyze:

- Compute a new ordering of the unknowns
- Compute the symbolic factorisation of the matrix
- Distribute column blocks and computation on processors

This function is not needed to use Murge interface, it only forces analyze step when user wants.

If this function is not used, analyze step will be performed when getting new distribution from MURGE, or filling the matrix.

MURGE_Factorize

```
INTS MURGE_Factorize ( INTS id);
```

```
SUBROUTINE MURGE_FACTORIZE ( ID, IERROR)
  INTS, INTENT(IN)      :: ID
  INTS, INTENT(OUT)     :: IERROR
END SUBROUTINE MURGE_FACTORIZE
```

Parameters :

id : Solver instance identification number.

Perform matrix factorization.

This function is not needed to use Murge interface, it only forces factorization when user wants.

If this function is not used, factorization will be performed with solve, when getting solution from MURGE.

MURGE_SetOrdering

```
INTS MURGE_SetOrdering ( INTS id, INTS * permutation);
```



```

SUBROUTINE MURGE_SETORDERING ( ID, PERMUTATION, IERROR)
    INTS, INTENT(IN)          :: ID
    INTS, INTENT(IN), DIMENSION(0) :: PERMUTATION
    INTS, INTENT(OUT)         :: IERROR
END SUBROUTINE MURGE_SETORDERING

```

Parameters :

id : Solver instance identification number.
permutation : Permutation to set internal computation ordering

Set permutation for PASTIX internal ordering.

The permutation array can be unallocated after the function is called.

MURGE_ForceNoFacto

```

INTS MURGE_ForceNoFacto ( INTS id);

SUBROUTINE MURGE_FORCENOFAC TO ( ID, IERROR)
    INTS, INTENT(IN)      :: ID
    INTS, INTENT(OUT)     :: IERROR
END SUBROUTINE MURGE_FORCENOFAC TO

```

Parameters :

id : Solver instance identification number.

Prevent Murge from running factorisation even if matrix has changed.

If an assembly is performed, next solve will trigger factorization except if this function is called between assembling the matrix and getting the solution.

MURGE_GetLocalProduct

```

INTS MURGE_GetLocalProduct ( INTS id, COEF * x);

SUBROUTINE MURGE_GETLOCALPRODUCT ( ID, x, IERROR)
    INTS, INTENT(IN)          :: ID
    COEF, INTENT(OUT), DIMENSION(0) :: X
    INTS, INTENT(OUT)         :: IERROR
END SUBROUTINE MURGE_GETLOCALPRODUCT

```

Parameters :

id : Solver instance identification number.
x : Array in which the local part of the product will be stored.

Perform the product $A \times x$ and returns its local part.

The vector must have been given through **MURGE_SetLocalRHS** or **MURGE_SetGlobalRHS**.

MURGE_GetGlobalProduct

```

INTS MURGE_GetGlobalProduct ( INTS id, COEF * x);

```

```

SUBROUTINE MURGE_GETGLOBALPRODUCT ( ID, x, IERROR)
  INTS, INTENT(IN)                :: ID
  COEF, INTENT(OUT), DIMENSION(0) :: X
  INTS, INTENT(OUT)                :: IERROR
END SUBROUTINE MURGE_GETGLOBALPRODUCT

```

Parameters :

- id** : Solver instance identification number.
- x** : Array in which the product will be stored.

Perform the product $A \times x$ and returns it globally.

The vector must have been given through MURGE_SetLocalRHS or MURGE_SetGlobalRHS.

MURGE_SetLocalNodeList

```

INTS MURGE_SetLocalNodeList ( INTS id,      INTS nodenbr
  (
                        INTS * nodelist);

SUBROUTINE MURGE_SETLOCALNODELIST ( ID, nodenbr, nodelist, IERROR)
  INTS, INTENT(IN)                :: ID
  INTS, INTENT(IN)                :: nodenbr
  INTS, INTENT(IN), DIMENSION(0) :: nodelist
  INTS, INTENT(OUT)                :: IERROR
END SUBROUTINE MURGE_SETLOCALNODELIST

```

Parameters :

- id** : Solver instance identification number.
- nodenbr** : Number of local nodes.
- nodelist** : Array containing global indexes of local nodes.

Set the distribution of the solver, preventing the solver from computing its own.

NEEDS TO BE CHECKED !

MURGE_AssemblySetSequence

```

INTS MURGE_AssemblySetSequence ( INTS id ,      INTL coefnbr,
                                INTS * ROWs, INTS * COLs,
                                INTS op,      INTS op2,
                                INTS mode, INTS nodes,
                                INTS * id_seq);

SUBROUTINE MURGE_ASSEMBLYSETSEQUENCE ( ID, coefnbr, ROWs, COLs,
                                op, op2, mode, nodes,
                                id_seq, IERROR)
  INTS, INTENT(IN)                :: ID
  INTL, INTENT(IN)                :: coefnbr
  INTS, INTENT(IN), DIMENSION(0) :: ROWs, COLs
  INTS, INTENT(IN)                :: op, op2, mode, nodes
  INTS, INTENT(OUT)                :: id_seq
  INTS, INTENT(OUT)                :: IERROR
END SUBROUTINE MURGE_ASSEMBLYSETSEQUENCE

```

Parameters :

- id** : Solver instance identification number.
- coefnbr** : Number of local entries in the sequence.
- ROWS** : List of rows of the sequence.
- COLs** : List of columns of the sequence.
- op** : Operation to perform for coefficient which appear several tim (see MURGE_ASSEMBLY_OP).
- op2** : Operation to perform when a coefficient is set by two different processors (see MURGE_ASSEMBLY_OP).
- mode** : Indicates if user ensure he will respect solvers distribution (see MURGE_ASSEMBLY_MODE).
- nodes** : Indicate if entries are given one by one or by node :
 - 0 : entries are entered value by value,
 - 1 : entries are entries node by node.

id_seq : Sequence ID.

Create a sequence of entries to build a matrix and store it for being reused.

MURGE_AssemblyUseSequence

```
INTS MURGE_AssemblyUseSequence ( INTS id , INTS id_seq,
                                COEF * values);

SUBROUTINE MURGE_ASSEMBLYUSESEQUENCE ( ID, id_seq, values, IERROR)
  INTS, INTENT(IN)          :: ID
  INTS, INTENT(IN)          :: id_seq
  COEF, INTENT(IN), DIMENSION(0) :: values
  INTS, INTENT(OUT)         :: IERROR
END SUBROUTINE MURGE_ASSEMBLYUSESEQUENCE
```

Parameters :

- id** : Solver instance identification number.
- id_seq** : Sequence ID.
- values** : Values to insert in the matrix.

Assembly the matrix using a stored sequence.

MURGE_AssemblyDeleteSequence

```
INTS MURGE_AssemblyDeleteSequence ( INTS id , INTS id_seq);

SUBROUTINE MURGE_ASSEMBLYDELETESEQUENCE ( ID, id_seq, IERROR)
  INTS, INTENT(IN)          :: ID
  INTS, INTENT(IN)          :: id_seq
  INTS, INTENT(OUT)         :: IERROR
END SUBROUTINE MURGE_ASSEMBLYDELETESEQUENCE
```

Parameters :

- id** : Solver instance identification number.
- id_seq** : Sequence ID.

Destroy an assembly sequence.

4.4 Auxiliary PaStiX functions

4.4.1 Distributed mode dedicated functions

Getting local nodes number

```
pastix_int_t pastix_getLocalNodeNbr ( pastix_data_t ** pastix_data );
```

```
SUBROUTINE PASTIX_FORTRAN_GETLOCALNODENBR ( PASTIX_DATA,  
                                              NODENBR)  
    pastix_data_ptr_t, INTENT(INOUT) :: PASTIX_DATA  
    INTS, INTENT(OUT) :: NODENBR  
END SUBROUTINE PASTIX_FORTRAN_GETLOCALNODENBR
```

Parameters :

pastix_data : Area used to store information between calls.

Return the node number in the new distribution computed by the analyze step
(Analyse step must have already been executed).

Getting local nodes list

```
int pastix_getLocalNodeLst ( pastix_data_t ** pastix_data,  
                             pastix_int_t * nodelst );
```

```
SUBROUTINE PASTIX_FORTRAN_GETLOCALNODELST ( PASTIX_DATA,  
                                              NODELST,  
                                              IERROR)  
    pastix_data_ptr_t, INTENT(INOUT) :: PASTIX_DATA  
    INTS, INTENT(OUT), DIMENSION(0) :: NODELST  
    INTS, INTENT(OUT) :: IERR  
END SUBROUTINE PASTIX_FORTRAN_GETLOCALNODELST
```

Parameters :

pastix_data : Area used to store information between calls.

nodelst : Array to receive the list of local nodes.

Fill **nodelst** with the list of local nodes

(**nodelst** must be at least **nodenbr*sizeof(pastix_int_t)**, where **nodenbr** is obtained from **pastix_getLocalNodeNbr**).

Getting local unknowns number

```
pastix_int_t pastix_getLocalUnknownNbr ( pastix_data_t ** pastix_data);
```

```
SUBROUTINE PASTIX_FORTRAN_GETLOCALUNKNOWNNBR ( PASTIX_DATA,  
                                                UNKNOWNNBR)  
    pastix_data_ptr_t, INTENT(INOUT) :: PASTIX_DATA  
    INTS, INTENT(OUT) :: UNKNOWNNBR  
END SUBROUTINE PASTIX_FORTRAN_GETLOCALUNKNOWNNBR
```

Parameters :

pastix_data : Area used to store information between calls.

Return the number of unknowns in the new distribution computed by the preprocessing.

Needs the preprocessing to be runned with **pastix_data** before.

Getting local unknowns list

```
int pastix_getLocalUnknownLst ( pastix_data_t ** pastix_data,
                                pastix_int_t  * unknownlst );

SUBROUTINE PASTIX_FORTRAN_GETLOCALUNKNOWNLST ( PASTIX_DATA,
                                                UNKNOWNLST,
                                                IERROR)

    pastix_data_ptr_t, INTENT(INOUT)      :: PASTIX_DATA
    INTS, INTENT(OUT), DIMENSION(0)      :: UNKNOWNLST
    INTS, INTENT(OUT)                     :: IERR
END SUBROUTINE PASTIX_FORTRAN_GETLOCALUNKNOWNLST
```

Parameters :

pastix_data : Area used to store information between calls.

nodelst : An array where to write the list of local nodes/columns.

Fill in **unknownlst** with the list of local unknowns/column.

Needs **unknownlst** to be allocated with **unknownnbr*sizeof(pastix_int_t)**, where **unknownnbr** has been computed by **pastix_getLocalUnknownNbr**.

4.4.2 Binding thread

```
void pastix_setBind ( pastix_data_t ** pastix_data, int thrdnbr,
                     int               * bindtab );

SUBROUTINE PASTIX_FORTRAN_SETBINDTAB ( PASTIX_DATA,
                                       THRDNBR,
                                       BINDTAB)

    pastix_data_ptr_t, INTENT(INOUT)      :: PASTIX_DATA
    INTS, INTENT(OUT)                     :: THRDNBR
    INTS, INTENT(OUT), DIMENSION(0)      :: BINDTAB
END SUBROUTINE PASTIX_FORTRAN_SETBINDTAB
```

Parameters :

pastix_data : Area used to store information between calls.

thrdnbr : Number of threads (== length of **bindtab**).

bindtab : List of processors for threads to be binded on.

Assign threads to processors.

Thread number **i** (starting with 0 in C and 1 in Fortran) will be binded to the core **bindtab[i]**

4.4.3 Working on CSC or CSCD

Checking and correcting the CSC or CSCD matrix

```
void pastix_checkMatrix ( MPI_Comm    pastix_comm, int      verb,
                        int          flagsym,   int      flagcor,
                        pastix_int_t n,         pastix_int_t ** colptr,
                        pastix_int_t ** row,     pastix_float_t ** avals,
                        pastix_int_t ** loc2glob ); int      dof
```

```
SUBROUTINE PASTIX_FORTRAN_CHECKMATRIX ( DATA_CHECK
                                         PASTIX_COMM,
                                         VERB,
                                         FLAGSYM,
                                         FLAGCOR,
                                         N,
                                         COLPTR,
                                         ROW,
                                         AVALS,
                                         LOC2GLOB)

    pastix_data_ptr_t, INTENT(OUT)      :: DATA_CHECK
    MPI_COMM, INTENT(IN)                :: PASTIX_COMM
    INTEGER, INTENT(IN)                 :: VERB
    INTEGER, INTENT(IN)                 :: FLAGSYM
    INTEGER, INTENT(IN)                 :: FLAGCOR
    pastix_int_t, INTENT(IN)            :: N
    pastix_int_t, INTENT(IN), DIMENSION(0) :: COLPTR
    pastix_int_t, INTENT(IN), DIMENSION(0) :: ROW
    pastix_int_t, INTENT(IN), DIMENSION(0) :: AVALS
    pastix_int_t, INTENT(IN), DIMENSION(0) :: LOC2GLOB
END SUBROUTINE PASTIX_FORTRAN_CHECKMATRIX
```

```
SUBROUTINE PASTIX_FORTRAN_CHECKMATRIX_END ( DATA_CHECK
                                             VERB,
                                             ROW,
                                             AVALS)

    pastix_data_ptr_t, INTENT(IN)      :: DATA_CHECK
    INTEGER, INTENT(IN)                :: VERB
    pastix_int_t, INTENT(IN), DIMENSION(0) :: ROW
    pastix_int_t, INTENT(IN), DIMENSION(0) :: AVALS
END SUBROUTINE PASTIX_FORTRAN_CHECKMATRIX_END
```

Parameters :

pastix_comm	: PASTIX MPI communicator.
verb	: Verbose mode (see Verbose modes).
flagsym	: Indicates if the matrix is symmetric (see Symmetric modes).
flagcor	: Indicates if the matrix can be reallocated (see Boolean modes).
n	: Matrix dimension.
colptr, row, avals	: Matrix in CSC format.
loc2glob	: Local to global column number correspondance.

Check and correct the user matrix in CSC format :

- Renumbers in Fortran numerotation (base 1) if needed (base 0)
- Can scale the matrix if compiled with `-DMC64 -DSCALING` (untested)
- Checks the symetry of the graph in non symmetric mode. With non distributed matrices, with `flagcor == APIYES`, tries to correct the matrix.

- sort the CSC.

In fortran, with correction enable, CSC array can be reallocated.

PASTIX works on a copy of the CSC and stores it internally if the number of entries changed.

If the number of entries changed ($colptr[n] - 1$), user as to reallocate rows and avals and then call PASTIX_FORTRAN_CHECKMATRIX_END().

Checking the symetry of a CSCD

```
int cscd_checksymb ( pastix_int_t    n,      pastix_int_t * ia,
                    pastix_int_t * ja,      pastix_int_t * l2g,
                    MPI_Comm    comm );
```

Parameters :

- n** : Number of local columns.
- ia** : Starting index of each column in **ja**.
- ja** : Row of each element.
- l2g** : Global column numbers of local columns.

Check the graph symmetry.

Correcting the symetry of a CSCD

```
int cscd_symgraph ( pastix_int_t    n,      pastix_int_t * ia,
                   pastix_int_t * ja,      pastix_float_t * a,
                   pastix_int_t * newn,    pastix_int_t ** newia,
                   pastix_int_t ** newja,  pastix_float_t ** newa,
                   pastix_int_t * l2g,    MPI_Comm    comm,
```

Parameters :

- n** : Number of local columns.
- ia** : Starting index of each column in **ja** and **a**.
- ja** : Row of each element.
- a** : Value of each element.
- newn** : New number of local columns.
- newia** : Starting index of each columns in **newja** and **newa**.
- newja** : Row of each element.
- newa** : Values of each element.
- l2g** : Global number of each local column.
- comm** : MPI communicator.

Symmetrize the graph.

Adding a CSCD into an other one

```
int cscd_addlocal ( pastix_int_t      n,      pastix_int_t * ia,
                   pastix_int_t      * ja,    pastix_float_t * a,
                   pastix_int_t      * l2g,    pastix_int_t   addn,
                   pastix_int_t      * addia, pastix_int_t   * addja,
                   pastix_float_t     * adda,    pastix_int_t   * addl2g,
                   pastix_int_t      * newn,    pastix_int_t   ** newia,
                   pastix_int_t      ** newja, pastix_float_t ** newa,
                   CSCD_OPERATIONS_t  OP );
```

Parameters :

n : Size of first CSCD matrix (same as newn).
ia : Column starting positions in first CSCD matrix.
ja : Rows in first CSCD matrix.
a : Values in first CSCD matrix (can be NULL).
l2g : Global column number map for first CSCD matrix.
addn : Size of the second CSCD matrix (to be added to base).
addia : Column starting positions in second CSCD matrix.
addja : Rows in second CSCD matrix.
adda : Values in second CSCD (can be NULL \rightarrow add \emptyset).
addl2g : Global column number map for second CSCD matrix.
newn : Size of output CSCD matrix (same as n).
newia : Column starting positions in output CSCD matrix.
newja : Rows in output CSCD matrix.
newa : Values in output CSCD matrix.
malloc_flag : Flag: Function call is internal to PASTIX.
OP : Specifies treatment of overlapping CSCD elements.

Add the second CSCD to the first CSCD, result is stored in the third CSCD (allocated in the function).

The operation OP can be :

CSCD_ADD : to add common coefficients.

CSCD_KEEP : to keep the coefficient of the first matrix.

CSCD_MAX : to keep the maximum value.

CSCD_MIN : to keep the minimum value.

CSCD_OVW : to overwrite with the value from the added CSCd.

Building a CSCD from a CSC

```
void csc_dispatch ( pastix_int_t      gN,      pastix_int_t * gcolptr,
                   pastix_int_t * grow,      pastix_float_t * gaval,
                   pastix_float_t * grhs,     pastix_int_t * gperm,
                   pastix_int_t * ginvp,
                   pastix_int_t * lN,      pastix_int_t ** lcolptr,
                   pastix_int_t ** lrow,    pastix_float_t ** laval,
                   pastix_float_t ** lrhs,   pastix_int_t ** lperm,
                   pastix_int_t ** loc2glob, int      dispatch,
                   MPI_Comm      pastix_comm );
```


Parameters :

gN	: Global CSC matrix number of columns.
gcolptr, grows, gaval	: Global CSC matrix
gperm	: Permutation table for global CSC matrix.
ginvp	: Inverse permutation table for global CSC matrix.
lN	: Local number of columns (output).
lcolptr, lrows, laval	: Local CSCD matrix (output).
lrhs	: Local part of the right hand side (output).
lperm	: Local part of the permutation table (output).
loc2glob	: Global numbers of local columns (before permutation).
dispatch	: Dispatching mode: CSC_DISP_SIMPLE Cut in n_{proc} parts of consecutive columns CSC_DISP_CYCLIC Use a cyclic distribution.
pastix_comm	: PaStiX MPI communicator.

Distribute a CSC into a CSCD.

In Fortran the routine as to be called in two steps, the first one compute the new CSCD and return its number of column and non-zeros, and the second one will copy the new CSCD into user's arrays.

```

SUBROUTINE CSC_DISPATCH_FORTRAN ( CSC_DATA, GN, GCOLPTR,
                                GROW, GAVALS,
                                GRHS, GPERM,
                                GINVP, DISPATCH,
                                NEWN, NEWNNZ,
                                PASTIX_COMM)
    pastix_data_ptr_t, INTENT(OUT)      :: CSC_DATA
    pastix_int_t, INTENT(IN)             :: GN
    pastix_int_t, INTENT(IN), DIMENSION(0) :: GCOLPTR
    pastix_int_t, INTENT(IN), DIMENSION(0) :: GROW
    pastix_int_t, INTENT(IN), DIMENSION(0) :: GAVALS
    pastix_int_t, INTENT(IN), DIMENSION(0) :: GRHS
    pastix_int_t, INTENT(IN), DIMENSION(0) :: GPERM
    pastix_int_t, INTENT(IN), DIMENSION(0) :: GINVP
    INTEGER, INTENT(IN)                  :: DISPATCH
    INTEGER, INTENT(OUT)                  :: NEWN
    INTEGER, INTENT(OUT)                  :: NEWNNZ
    MPL_COMM, INTENT(IN)                  :: PASTIX_COMM
END SUBROUTINE CSC_DISPATCH_FORTRAN

```

```

SUBROUTINE CSC_DISPATCH_FORTRAN_END ( CSC_DATA, LCOLPTR,
                                    LROW, LAVALS,
                                    LRHS, LPERM,
                                    L2G)
    pastix_data_ptr_t, INTENT(OUT)      :: CSC_DATA
    pastix_int_t, INTENT(IN), DIMENSION(n) :: LCOLPTR
    pastix_int_t, INTENT(IN), DIMENSION(nnz) :: LROW
    pastix_int_t, INTENT(IN), DIMENSION(nnz) :: LAVALS
    pastix_int_t, INTENT(IN), DIMENSION(n) :: LRHS
    pastix_int_t, INTENT(IN), DIMENSION(n) :: LPERM
    pastix_int_t, INTENT(IN), DIMENSION(n) :: L2G
END SUBROUTINE CSC_DISPATCH_FORTRAN_END

```

Changing a CSCD distribution

```
int cscd_redispatch ( pastix_int_t    n,      pastix_int_t * ia,
                     pastix_int_t * ja,      pastix_float_t * a,
                     pastix_float_t * rhs,    pastix_int_t * l2g,
                     pastix_int_t  dn,      pastix_int_t ** dia,
                     pastix_int_t ** dja,    pastix_float_t ** da,
                     pastix_float_t ** drhs, pastix_int_t * dl2g,
                     MPI_Comm    comm);
```

Parameters :

n : Number of local columns
ia : First cscd starting index of each column in **ja** and **a**
ja : Row of each element in first CSCD
a : Value of each CSCD in first CSCD (can be NULL)
rhs : Right-hand-side member corresponding to the first CSCD (can be NULL)
l2g : Local to global column numbers for first CSCD
dn : Number of local columns
dia : New CSCD starting index of each column in **ja** and **a**
dja : Row of each element in new CSCD
da : Value of each CSCD in new CSCD
rhs : Right-hand-side member corresponding to the new CSCD
dl2g : Local to global column numbers for new CSCD
comm : MPI communicator

Redistribute the first cscd, distributed with **l2g** local to global array, into a new one using **dl2g** as local to global array.

The algorithm works in four main steps :

- gather all new loc2globs on all processors;
- allocate **dia**, **dja** and **da**;
- Create new CSC for each processor and send it;
- Merge all new CSC to the new local CSC with **cscd_addlocal()**.

If communicator size is one, check that $n = dn$ and $l2g = dl2g$ and simply create a copy of the first CSCD.

In Fortran the function as to be called in to step, the first one, **CSCD_REDISPATCH_FORTRAN**, to compute the new CSCD, and the second one, **CSCD_REDISPATCH_FORTRAN_END** to copy the computed CSCD into the user allocated structure.

```
SUBROUTINE CSCD_REDISPATCH_FORTRAN ( CSC_DATA, N, IA
                                     JA, A,
                                     RHS, L2G,
                                     NEWN, NEWL2G,
                                     FORTRAN_COMM, IERR)
```

```

    pastix_data_ptr_t, INTENT(OUT)      :: CSC_DATA
    pastix_int_t, INTENT(IN)             :: N
    pastix_int_t, INTENT(IN), DIMENSION(0) :: IA
    pastix_int_t, INTENT(IN), DIMENSION(0) :: JA
    pastix_int_t, INTENT(IN), DIMENSION(0) :: A
    pastix_int_t, INTENT(IN), DIMENSION(0) :: RHS
    pastix_int_t, INTENT(IN), DIMENSION(0) :: L2G
    pastix_int_t, INTENT(IN)             :: NEWN
    pastix_int_t, INTENT(IN), DIMENSION(0) :: NEWL2G
    pastix_int_t, INTENT(OUT)             :: NEWNNZ
    MPLCOMM, INTENT(IN)                   :: FORTRAN_COMM
    INTEGER, INTENT(OUT)                   :: IERR
END SUBROUTINE CSCD_REDISPATCH_FORTRAN

SUBROUTINE CSCD_REDISPATCH_FORTRAN_END ( CSC_DATA, DCOLPTR
                                         DROW, DAVALS,
                                         DRHS)
    pastix_data_ptr_t, INTENT(OUT)      :: CSC_DATA
    pastix_int_t, INTENT(IN), DIMENSION(n) :: LCOLPTR
    pastix_int_t, INTENT(IN), DIMENSION(nnz) :: LROW
    pastix_int_t, INTENT(IN), DIMENSION(nnz) :: LAVALS
    pastix_int_t, INTENT(IN), DIMENSION(n) :: LRHS
END SUBROUTINE CSCD_REDISPATCH_FORTRAN_END

```

4.4.4 Schur complement

Schur can be obtained through two ways :

- User can set is unknown list and get a copy of the schur.
- User can set is unknown list, ask for the schur distribution and give a memory area in which PASTIX will store the schur.

This second option permit to optimize memory consumption.

Indicating schur complement indices

```

int pastix_setSchurUnknownList ( pastix_data_t * pastix_data,
                                pastix_int_t   n,
                                pastix_int_t   * list );

```

Parameters :

pastix_data : Area used to store information between calls.
n : Number of unknowns.
list : List of unknowns.

Returns :

NO_ERR : If all goes well.

Set the list of unknowns composing the schur complement.

This function must be used with **IPARM_SCHUR** set to **API_YES**.

This function must be called before the graph partitioning step.

After using it and performing factorization, the **schur** complement can be obtained with **pastix_getSchur** (4.4.4) or **pastix_setSchurArray** (4.4.4) and the following solve will be performed on the non-schur part of the matrix (but using a full lenght right-hand-side).

```

SUBROUTINE PASTIX_FORTRAN_SETCHURUNKNOWNLIST ( PASTIX_DATA,
                                                N, IA
                                                LIST)

    pastix_data_ptr_t, INTENT(OUT)      :: PASTIX_DATA
    pastix_int_t, INTENT(IN)            :: N
    pastix_int_t, INTENT(IN), DIMENSION(0) :: LIST
END SUBROUTINE PASTIX_FORTRAN_SETCHURUNKNOWNLIST

```

Getting a copy the schur complement

```

int pastix_getSchur ( pastix_data_t * pastix_data,
                     pastix_float_t * schur );

```

Parameters :

pastix_data : Area used to store information between calls.
schur : Array to fill-in with Schur complement.

Returns :

NO_ERR : If all goes well.

Fill the array **schur** with the schur complement. The **schur** array must be allocated with n^2 **pastix_float_t**, where **n** is the number of unknowns in the schur complement.

```

SUBROUTINE PASTIX_FORTRAN_GETSCHUR ( PASTIX_DATA,
                                      SCHUR)

    pastix_data_ptr_t, INTENT(OUT) :: PASTIX_DATA
    pastix_float_t, INTENT(INOUT)  :: SCHUR
END SUBROUTINE PASTIX_FORTRAN_GETSCHUR

```

Getting schur distribution

```

int pastix_getSchurLocalNodeNbr ( pastix_data_t * pastix_data,
                                 pastix_int_t   * schurLocalNodeNbr );

```

Parameters :

pastix_data : Area used to store information between calls.
schurLocalNodeNbr : Number of nodes in local part of the schur.

Returns :

NO_ERR : If all goes well.

Get the number of nodes in the part of the **schur** local to the MPI task.

With this information, user can allocate the local part of the **schur** complement and give to PASTIX with **pastix_setSchurArray** (4.4.4)

This function must be called after the analysis and before the numerical factorization.

```

SUBROUTINE PASTIX_FORTRAN_GETSCHURLocalNODENBR ( PASTIX_DATA,
                                                  NODENBR, IERR)

    pastix_data_ptr_t, INTENT(INOUT) :: PASTIX_DATA
    pastix_int_t, INTENT(OUT)        :: NODENBR , IERR
END SUBROUTINE PASTIX_FORTRAN_GETLOCALNODENBR

```

```

int pastix_getSchurLocalNodeList ( pastix_data_t * pastix_data,
                                 pastix_int_t   * schurLocalNodeList );

```

Parameters :

`pastix_data` : Area used to store information between calls.
`schurLocalNodeList` : List of the nodes in local part of the schur.

Returns :

`NO_ERR` : If all goes well.

Return the list of nodes of the schur complement which will be stored on the current MPI process. This knowledge is necessary for the user to interpret the part of the schur stored in the memory allocation he provides with `pastix_setSchurArray` (4.4.4).

This function can't be called before analysis step and user must have called `pastix_getSchurLocalNodeNbr` (4.4.4) to be aware of the size of the node list.

```
SUBROUTINE PASTIX_FORTRAN.GETSCHURLOCALNODELIST ( PASTIX_DATA,  
                                                    NODELIST, IERR)  
    pastix_data_ptr_t, INTENT(INOUT)      :: PASTIX_DATA  
    pastix_int_t, INTENT(OUT), DIMENSION(0) :: NODELIST  
    pastix_int_t, INTENT(OUT)              :: IERR  
END SUBROUTINE PASTIX_FORTRAN.GETLOCALNODELIST
```

—

```
int pastix_getSchurLocalUnkownNbr ( pastix_data_t * pastix_data,  
                                   pastix_int_t   * schurLocalUnkownNbr );
```

Parameters :

`pastix_data` : Area used to store information between calls.
`schurLocalUnkownNbr` : Number of unkowns in local part of the schur.

Returns :

`NO_ERR` : If all goes well.

Get the number of unkowns in the part of the `schur` local to the MPI task.

With this information, user can allocate the local part of the `schur` complement and give to PASTIX with `pastix_setSchurArray` (4.4.4)

This function must be called after the analysis and before the numerical factorization.

```
SUBROUTINE PASTIX_FORTRAN.GETSCHURLOCALUNKOWNNBR ( PASTIX_DATA,  
                                                    UNKOWNNBR, IERR)  
    pastix_data_ptr_t, INTENT(INOUT) :: PASTIX_DATA  
    pastix_int_t, INTENT(OUT)         :: UNKOWNNBR, IERR  
END SUBROUTINE PASTIX_FORTRAN.GETLOCALUNKOWNNBR
```

—

```
int pastix_getSchurLocalUnkownList ( pastix_data_t * pastix_data,  
                                   pastix_int_t   * schurLocalUnkownList );
```

Parameters :

`pastix_data` : Area used to store information between calls.
`schurLocalUnkownList` : List of the unkowns in local part of the schur.

Returns :

`NO_ERR` : If all goes well.

Return the list of unkowns of the schur complement which will be stored on the current MPI process.

This knowledge is necessary for the user to interpret the part of the schur stored in the memory allocation he provides with `pastix_setSchurArray` (4.4.4).

This function can't be called before analysis step and user must have called `pastix_getSchurLocalUnkownNbr` (4.4.4) to be aware of the size of the unkown list.

```

SUBROUTINE PASTIX_FORTRAN.GETSCHURLOCALUNKOWNLIST ( PASTIX_DATA,
                                                    UNKOWNLIST, IERR)

    pastix_data_ptr_t, INTENT(INOUT)                :: PASTIX_DATA
    pastix_int_t, INTENT(OUT), DIMENSION(0)         :: UNKOWNLIST
    pastix_int_t, INTENT(OUT)                       :: IERR
END SUBROUTINE PASTIX_FORTRAN.GETLOCALUNKOWNLIST

```

Setting memory allocation to store the schur

```

int pastix_setSchurArray ( pastix_data_t * pastix_data,
                           pastix_float_t * schurArray );

```

Parameters :

`pastix_data` : Area used to store information between calls.
`schurLocalUnkownList` : Memory area allocated by the user to store the local part of the schur complement.

Returns :

`NO_ERR` : If all goes well.

Using this fonction, the user provides a memory area for the storage of the schur.

The memory area must be of size $nSchurLocalcol \times nSchurCol$, where `nSchurLocalcol` is the number of local columns required for the MPI node and `nSchurCol` is the global number of columns of the `schur` (and row since the schur is a square).

This function as to be called after analysis step and before numerical factorization is performed. After the factorization, the `schur` column `schurLocalColumn[i]` will be stored from `schurArray[i × nSchurCol]` to `schurArray[(i + 1) × nSchurCol - 1]`.

```

SUBROUTINE PASTIX_FORTRAN.SETSCHURARRAY ( PASTIX_DATA,
                                           SCHURARRAY, IERR)

    pastix_data_ptr_t, INTENT(INOUT)                :: PASTIX_DATA
    PASTIX_FLOAT, INTENT(OUT), DIMENSION(0)         :: SCHURARRAY
    pastix_int_t, INTENT(OUT)                       :: IERR
END SUBROUTINE PASTIX_FORTRAN.SETSCHURARRAY

```

4.5 Multi-arithmetic

All PASTIX functions, except Murge's ones, are available with 5 differents arithmetics.

The default arithmetic, defined in `config.in` file with the flag `CCTYPESFLT`. All the functions corresponding to this arithmetic are listed above. For example, `pastix` (Fig. 4.1, p.19).

The simple float arithmetic, corresponding to all functions presented above prefixed by “s-”. For example, `s_pastix` (Fig. 4.8, p.38).

The double float arithmetic, prefixed by “d-”. For example, `d_pastix` (Fig. 4.9, p.38).

The simple complex arithmetic, prefixed by “c-”. For example, `c_pastix` (Fig. 4.10, p.38).

The double complex arithmetic, prefixed by “z-”. For example, `z_pastix` (Fig. 4.11, p.38).

```

#include "pastix.h"
void s_pastix ( pastix_datat_t ** pastix_data, MPI_Comm    pastix_comm,
                pastix_int_t    n,                pastix_int_t * colptr,
                pastix_int_t    * row,            float      * avals,
                pastix_int_t    * perm,           pastix_int_t * invp,
                float            * b,              pastix_int_t rhs,
                pastix_int_t    * iparm,          double      * dparm );

```

Figure 4.8: PASTIX main function prototype, simple float mode

```

#include "pastix.h"
void d_pastix ( pastix_data_t ** pastix_data, MPI_Comm    pastix_comm,
                pastix_int_t    n,                pastix_int_t * colptr,
                pastix_int_t    * row,            double      * avals,
                pastix_int_t    * perm,           pastix_int_t * invp,
                double           * b,              pastix_int_t rhs,
                pastix_int_t    * iparm,          double      * dparm );

```

Figure 4.9: PASTIX main function prototype, double float mode

```

#include "pastix.h"
void c_pastix ( pastix_data_t ** pastix_data, MPI_Comm    pastix_comm,
                pastix_int_t    n,                pastix_int_t * colptr,
                pastix_int_t    * row,            float complex * avals,
                pastix_int_t    * perm,           pastix_int_t * invp,
                float complex    * b,              pastix_int_t rhs,
                pastix_int_t    * iparm,          double      * dparm );

```

Figure 4.10: PASTIX main function prototype, complex mode

```

#include "pastix.h"
void z_pastix ( pastix_datat_t ** pastix_data, MPI_Comm    pastix_comm,
                pastix_int_t    n,                pastix_int_t * colptr,
                pastix_int_t    * row,            double complex * avals,
                pastix_int_t    * perm,           pastix_int_t * invp,
                double complex    * b,              pastix_int_t rhs,
                pastix_int_t    * iparm,          double      * dparm );

```

Figure 4.11: PASTIX main function prototype, double complex mode

Chapter 5

SWIG python wrapper

A python wrapper has been added to PASTIX using SWIG. It gives access to the whole original sequential and distributed interface.

The MURGE interface is not available yet for python.

It should be possible to had wrapper to all available SWIG outputs.

5.1 Requirement

To build the python interface, you need SWIG to generate the interface and MPI4PY to be able to use MPI in python.

5.2 Building the python wrapper

SWIG python wrapper requires librairies compiled with the `position-independent code` option (`-fPIC` in `gcc`, `gfortran`, `icc` or `ifort`).

Thus, MPI, SCOTCH (or METIS) and PASTIX must be compiled with `-fPIC`

To compile PaStiX with this option and to compile the python wrapper, uncomment the corresponding section of the `config.in` file.

NB : All used libraries must also be build with the `position-independent` option.

5.3 Example

An example using PASTIX in python as been written in `example/src/pastix.python.c`.

You can call it by running `make python` or the command `PYTHONPATH=$(MPI4PY_LIBDIR):$(PASTIX_LIBDIR) python example/src/pastix.python.py`.

Chapter 6

Examples

Many different examples are provided with PASTIX library sources.

These examples are meant to be simple and documented to simplify user's work when including PASTIX in their softwares.

Examples are stored in `src/example/src/` and compiled with `make examples`. The resulting binaries are stored in `src/example/bin/` directory.

6.1 Examples in C

As PASTIX is written in C, there are many examples in this language.

6.1.1 Parameters

All this examples are sharing the same parameter option that can be listed with `example/bin/simple` or any other example without parameters or with `-h`.

Usage : `./example/bin/simple [option]`

options :

<code>-rsa</code>	<code>[filename]</code>	driver RSA (use Fortran)
<code>-chb</code>	<code>[filename]</code>	driver CHB
<code>-ccc</code>	<code>[filename]</code>	driver CCC
<code>-rcc</code>	<code>[filename]</code>	driver RCC
<code>-olaf</code>	<code>[filename]</code>	driver OLAF
<code>-peer</code>	<code>[filename]</code>	driver PEER
<code>-hb</code>	<code>[filename]</code>	driver HB (double)
<code>-3files</code>	<code>[filename]</code>	driver IJV 3files
<code>-mm</code>	<code>[filename]</code>	driver Matrix Market
<code>-ord</code>	<code><scotch metis></code>	select ordering library
<code>-lap</code>	<code><integer></code>	generate a laplacian of size <code><integer></code>
<code>-incomp</code>	<code><integer> <integer></code>	incomplete factorization, with the given level of fill [1-5] and amalgamation [10-70]
<code>-ooc</code>	<code><integer></code>	Memory limit in Mo/percent depending on compilation options
<code>-kass</code>	<code><integer></code>	kass, with the given amalgamation

-t	<integer>	define thread number
-v	<integer>	define verbose level (1,2 or 3)
-h		print this help

6.1.2 simple.c

The `simple.c` example is an example simply using PASTIX with no particular option. It reads the matrix, check the matrix structure, set options depending on the matrix (ie: `IPARM_SYM`), run PASTIX on the builded system. After that the returned solution is tested.

The compilation of the example generate `simple`, `ssimple`, `dsimple`, `csimple` and `zsimple`, each one corresponding to one arithmetic.

This example is tested every day with our regression tests.

6.1.3 simple_dist.c

This example is similar to `simple.c` example, except that it uses the distributed interface. It read and check the matrix, set the parameters, run preprocessing steps. After that it get solvers distribution, redistribute the matrix to follow it and run the factorization and solve steps.

This example also produce five executables corresponding to the five arithmetics.

This example is also tested with our daily runs.

6.1.4 do_not_redispatch_rhs.c

In this example the factorization is computed using user's matrix in it's original distribution, then the right-hand-side member is built in PASTIX distribution.

To allow this kind of mixed usage, user has to set `IPARM_RHSD_CHECK` to `API_NO`.

6.1.5 step-by-step.c

This examples call each PASTIX steps one after one. `nfact` - which is 2 in the code - factorisation are called, and, for each one, `nsolv` - 2 - solving steps are performed.

6.1.6 step-by-step_dist.c

This example is similar to `step-by-step` one but calls are performed with distributed matrices. Pre-processing is performed with user's distribution, then the matrices are re-distributed in the solver's distribution to perform several factorizations and resolutions.

6.1.7 reentrant.c

In this example, PASTIX is called from several threads. This example is here to check that the library can be called from any number of threads, which could be problematic.

6.1.8 multi-comm.c

This example check the possibility to call PASTIX in different MPI communicators.

6.1.9 multi-comm-step.c

This example is similar to multi-comm.c but it also uses step by step calls, as in step-by-step.c

6.1.10 schur.c

In this example, we build the Schur complement of the matrix. The user provide a list of unknown corresponding to the Schur complement, run the matrix factorization and then get the Schur complement as a dense block array.

After that we can perform updown step on the non-schur part of the matrix.

6.1.11 isolate_zeros.c

In this example, PASTIX first isolate zeros diagonal terms at the end of the matrix before factorizing the matrix.

6.2 Examples in Fortran

As many PASTIX users are Fortran users, we provide some examples using our fortran interface to PASTIX

6.2.1 Parameters

All this examples are sharing the same parameter option that can be listed with `example/bin/simple -h`.

All this examples - except the murge one - are sharing the same parameter option that can be listed with `example/bin/fsimple` or any other example without parameters or with `-h`.

Usage : `./example/bin/fsimple [option]`

options :

<code>-rsa</code>	<code>[filename]</code>	default driver RSA (use Fortran)
<code>-lap</code>	<code><integer></code>	generate a laplacian of size <code><integer></code>
<code>-t</code>	<code><integer></code>	define thread number
<code>-v</code>	<code><integer></code>	verbose level (1, 2 or 3)
<code>-h</code>		print this help

6.2.2 fsimple.f90

This example correspond to `simple.c` example, but it is written in Fortran.

6.2.3 fstep-by-step.f90

This example is the Fortran version of `step-by-step.c`

6.2.4 Murge-Fortran.f90

In this example we use MURGE interface to solve a system in **Fortran**.

The system correspond to a Laplacian.

Chapter 7

Outputs

7.1 Logs

This section aim at giving clue to analyse PASTIX outputs on standard output and error streams.

7.1.1 Controlling output logs

The parameter `IPARM_VERBOSE` can be use to modify output logs.

The different values than can be given are :

`API_VERBOSE_NOT` : No output at all.

`API_VERBOSE_NO` : Few output lines.

`API_VERBOSE_YES` : Many output lines.

`API_VERBOSE_CHATTERBOX` : Like a real gossip.

`API_VERBOSE_UNBEARABLE` : Please stop it ! It's talking too much !

7.1.2 Interpreting output logs

% Log text		: Signification
% <code>Version</code>	:	: Subversion's version number.
% <code>SMP_SOPALIN</code>	:	: Indicate if <code>SMP_SOPALIN</code> as been defined.
% <code>VERSION_MPI</code>	:	: Indicate if <code>NO_MPI</code> as been defined.
% <code>PASTIX_DYNSCHED</code>	:	: Indicate if <code>PASTIX_DYNSCHED</code> as been defined.
% <code>STATS_SOPALIN</code>	:	: Indicate if <code>STATS_SOPALIN</code> as been defined.
% <code>NAPA_SOPALIN</code>	:	: Indicate if <code>NAPA_SOPALIN</code> as been defined.
% <code>TEST_IRecv</code>	:	: Indicate if <code>TEST_IRecv</code> as been defined.
% <code>TEST_ISEND</code>	:	: Indicate if <code>TEST_ISEND</code> as been defined.
% <code>THREAD_COMM</code>	:	: Indicate if <code>THREAD_COMM</code> as been defined.
% <code>THREAD_FUNNELED</code>	:	: Indicate if <code>THREAD_FUNNELED</code> as been defined.
% <code>TAG</code>	:	: Indicate the communication tag mode.
% <code>OUT_OF_CORE</code>	:	: Indicate if <code>OOC</code> as been defined.
% <code>DISTRIBUTED</code>	:	: Indicate if <code>DISTRIBUTED</code> as been defined.
% <code>FORCE_CONSO</code>	:	: Indicate if <code>FORCE_CONSO</code> as been defined.

continued on next page ...

...continued@from previous page

```
% RECV_FANIN_OR_BLOCK      :      : Indicate if RECV_FANIN_OR_BLOCK as been defined.
% FLUIDBOX                  :      : Indicate if FLUIDBOX as been defined.
% METIS                     :      : Indicate if METIS as been defined.
% INTEGER TYPE              :      : Indicate integer type.
% FLOAT TYPE                :      : Indicate float type.
% Ordering                  :      : Starting ordering step.
%   > Symmetrizing graph    :      : Symmetrizing the graph.
%   > Removing diag         :      : Removing diagonal entries.
%   > Initiating ordering   :      : Starting ordering routines.
% Symbolic Factorization :      : Starting symbolic factorization.
% Kass                      :      : Special step needed if not using Scotch.
% Analyse                  :      : Starting analyse and distribution step.
% Numerical Factorization :      : Starting numerical factorization.
% Solve                   :      : Starting up-down step.
% Reffinement              :      : Starting refinement step.
%   Time to compute ordering :      : Time spend to perform ordering.
%   Number of cluster       :      : Number of MPI process.
%   Number of processor per cluster :      : Total number of threads.
%   Number of thread number per MPI process :      : Explicit, isn't it ?
%   Check the symbol matrix :      : Checking symbolic matrix given for analyse step.
%   Check the solver structure :      : Checking solver structure built for numerical steps (factorisation, up-down, refinement).
%
%   Building elimination graph :      : All following are analyse steps...
%   Re-Building elimination graph :      : ...
%   Building cost matrix        :      : ...
%   Building elimination tree   :      : ...
%   Building task graph         :      : ...
%   Number of tasks             :      : Number of tasks computed by analyse step.
%   Distributing partition      :      : Distributing step.
%   Time to analyze             :      : Analyse step.
%   Number of nonzeros in factorized matrice :      : Number of non zeros in factorized matrix computed by analyse step.
%
% P : Number of nonzeros (local block structure) :      : Number of non zeros on processor P
% P : SolverMatrix size (without coefficients)    :      : Size of the structure used for numerical steps on processor P.
% Fill-in                                         :      : Theoretical fill-in of the matrix
% Number of operations (LU)                     :      : Number of operation for factorization.
% Number of operations (LLt)                    :      : Number of operation for factorization.
% Prediction Time to factorize (IBM PWR5 ESSL)   :      : Prediction time based on IBM Power5 with Blas ESSL.
% Maximum coeftab size (cefficients)             :      : Maximum memory used to store values of the matrix used during factorization.
%
% --- Sopalin : Allocation de la structure globale --- :      : Allocating factorization structure.
% --- Fin Sopalin Init                                :      : End of factorization initialisation step.
% --- Initialisation des tableaux globaux             :      : Initialization of global arrays.
% --- Sopalin : Local structure allocation            :      : Allocation of local thread structures.
% --- Sopalin : Threads are NOT binded                :      : Explicit...
% --- Sopalin : Threads are binded                   :      : Explicit...
%   - P : Envois X - Receptions Y                    :      : Processor P will send X buffers and receive Y buffers.
% --- Sopalin Begin                                  :      : Starting factorization.
% --- Sopalin End                                    :      : End of factorization.
% --- Down Step                                       :      : Explicit...
% --- Diag Step                                       :      : Explicit...
% --- Up Step                                         :      : Explicit...
% Generate RHS for X=1                               :      : Generata a right-hand-side member such that  $\forall i, X_i = 1$ .
% Generate RHS for X=i                               :      : Generata a right-hand-side member such that  $\forall i, X_i = i$ .
% OOC memory limit                                   :      : Limit set for out-of-core mode.
% [P] IN STEP                                         :      : Title for out-of-core outputs on processor P.
% [P]   written X, allocated : Y                     :      : Amount of data written (X) and allocated (Y) on processor P during step STEP.
```

continued on next page ...

...continued from previous page

% [P] read	: Amount of data read on disk on processor P during step STEP.
% [P] Allocated	: Amount of data allocated on processor P at the end of step STEP.
% [P] Maximum allocated	: Maximum allocated data during until end of step STEP.
% - iteration N :	: Iteration number in refinement.
% time to solve	: Time to perform up-down during refinement step.
% total iteration time	: Total time of the iteration step.
% error	: Precision after iteration.
% r	: Rsidual error (ie. $Ax-b$) after iteration.
% b	: Right-hand-side member's norm.
% r / b	: Precision after iteration.
% Time to fill internal csc	: Time to fill internal structure from user data.
% Max memory used after factorization	: Maximum memory peak (reduced on MPI_MAX).
% Memory used after factorization	: Maximum current memory usage after factorization.
% Max memory used after clean	: Maximum memory peak (reduced on MPI_MAX).
% Memory used after clean	: Maximum current memory usage after clean (should be 0).
% Static pivoting	: Number of static pivoting performed.
% Inertia	: Inertia of the matrix (Number of non zeros on the diagonal)
% Number of tasks added by esp	: Number of task added by Enhanced Sparse Parallelism option.
% Time to factorize	: Maximum time to perform numrical factorization.
% Time to solve	: Maximum time to perform up-down.
% Refinement X iterations, norm=N	: Number of iterations during refinement and final precision.
% Time for refinement	: Maximum time to perform refinement.

Table 7.1: Interpreting output table

7.2 Integer and double outputs

This section present integer and double outputs parameters.

7.2.1 Integer parameters

IPARM_NBITER : Number of iteration during refinement.

IPARM_STATIC_PIVOTING : Number of static pivoting during factorization.

IPARM_NNZERO : Number of non-zeros computed during analyse.

IPARM_ALLOCATED_TERMS : Number of terms allocated for the matrix factorization (matrix storage and communication buffer's, sum over MPI process).

IPARM_NNZEROS_BLOCK_LOCAL : Number of non zeros allocated in the local matrix.

IPARM_INERTIA : Inertia of the matrix (Number of non zeros on the diagonal).

IPARM_ESP_NBTASKS : Number of tasks added by Enhanced Sparse Parallelism option.

IPARM_ERROR_NUMBER : Error number returned (see ERR_NUMBERS).

7.2.2 Double parameters

DPARM_FILL_IN : Fill-in ratio.

DPARM_MEM_MAX : Max memory used, during all execution, between all MPI process.

DPARM_RELATIVE_ERROR : Precision of the given solution.

DPARM_ANALYZE_TIME : Time to perform analyze.

DPARM_PRED_FACT_TIME : Predicted time to perform numerical factorization.

DPARM_FACT_TIME : Time to perform numerical factorization.

DPARM_SOLV_TIME : Time to compute up-down.

DPARM_FACT_FLOPS : Number of Floating Point operations per seconds during numerical factorization.

DPARM_SOLV_FLOPS : Number of Floating Point operations per seconds during up-down.

DPARM_RAFF_TIME : Refinement computation time.

Nomenclature

BLAS : Basic Linear Algebra Subprograms, de facto application programming interface standard for publishing libraries to perform basic linear algebra operations such as vector and matrix multiplication. They were first published in 1979, and are used to build larger packages such as LAPACK. Heavily used in high-performance computing, highly optimized implementations of the BLAS interface have been developed by hardware vendors such as by Intel as well as by other authors (e.g. ATLAS is a portable self-optimizing BLAS). The LINPACK benchmark relies heavily on DGEMM, a BLAS subroutine, for its performance. (en.wikipedia.org)

Cholesky decomposition : In mathematics, the Cholesky decomposition is named after André-Louis Cholesky, who found that a symmetric positive-definite matrix can be decomposed into a lower triangular matrix and the transpose of the lower triangular matrix. The lower triangular matrix is the Cholesky triangle of the original, positive-definite matrix. Cholesky's result has since been extended to matrices with complex entries. (en.wikipedia.org)

GMRES : In mathematics, the generalized minimal residual method is an iterative method for the numerical solution of a system of linear equations. The method approximates the solution by the vector in a Krylov subspace with minimal residual. The Arnoldi iteration is used to find this vector.

NUMA : Non-Uniform Memory Access or Non-Uniform Memory Architecture (NUMA) is a computer memory design used in multiprocessors, where the memory access time depends on the memory location relative to a processor. Under NUMA, a processor can access its own local memory faster than non-local memory, that is, memory local to another processor or memory shared between processors.

Separator : A set of vertices that connects two disjoint parts of the graph.

Index

Compilation options

FORCE_NOMPI, [10](#)
FORCE_NOSMP, [10](#)
INTSIZE32, [10](#)
INTSIZE64, [10](#)
MEMORY_USAGE, [11](#)
NO_MPI_TYPE, [11](#)
NUMA_ALLOC, [11](#)
OOC, [11](#)
PASTIX_USE_BUBBLE, [11](#)
THREAD_COMM, [11](#)
THREAD_MULTIPLE, [11](#)

config.in, [9](#), [15](#)

Double parameters

DPARM_ANALYZE_TIME, [16](#), [46](#)
DPARM_EPSILON_MAGN_CTRL, [17](#)
DPARM_EPSILON_REFINEMENT, [18](#)
DPARM_FACT_FLOPS, [47](#)
DPARM_FACT_TIME, [17](#), [47](#)
DPARM_FILL_IN, [46](#)
DPARM_MEM_MAX, [11](#), [46](#)
DPARM_PRED_FACT_TIME, [16](#), [46](#)
DPARM_RAFF_TIME, [18](#), [47](#)
DPARM_RELATIVE_ERROR, [18](#), [46](#)
DPARM_SOLV_FLOPS, [17](#), [47](#)
DPARM_SOLV_TIME, [17](#), [47](#)

Fonctions

CSCD_REDISPATCH_FORTRAN_END, [34](#)
CSCD_REDISPATCH_FORTRAN, [34](#)
CSC_DISPATCH_FORTRAN_END, [32](#)
CSC_DISPATCH_FORTRAN, [32](#)
MURGE_ANALYZE, [23](#)
MURGE_ASSEMBLYDELETSEQUENCE, [26](#)
MURGE_ASSEMBLYSETSEQUENCE, [25](#)
MURGE_ASSEMBLYUSESEQUENCE, [26](#)
MURGE_Analyze, [23](#)
MURGE_AssemblyDeleteSequence, [26](#)
MURGE_AssemblySetSequence, [25](#)
MURGE_AssemblyUseSequence, [26](#)
MURGE_FACTORIZE, [23](#)
MURGE_FORCENOFACCTO, [24](#)
MURGE_Factorize, [23](#)
MURGE_ForceNoFacto, [24](#)

MURGE_GETGLOBALPRODUCT, [25](#)
MURGE_GETLOCALPRODUCT, [24](#)
MURGE_GetGlobalProduct, [24](#)
MURGE_GetLocalProduct, [24](#)
MURGE_SETLOCALNODELIST, [25](#)
MURGE_SETORDERING, [24](#)
MURGE_SetLocalNodeList, [25](#)
MURGE_SetOrdering, [23](#)
PASTIX_FORTRAN_CHECKMATRIX_END, [29](#)
PASTIX_FORTRAN_CHECKMATRIX, [29](#)
PASTIX_FORTRAN_GETLOCALNODELST, [27](#)
PASTIX_FORTRAN_GETLOCALNODENBR, [27](#)
PASTIX_FORTRAN_GETLOCALUNKNOWNLST, [28](#)
PASTIX_FORTRAN_GETLOCALUNKNOWNNBR, [27](#)
PASTIX_FORTRAN_GETSCHURLOCALNODELIST, [36](#)
PASTIX_FORTRAN_GETSCHURLOCALNODENBR, [35](#)
PASTIX_FORTRAN_GETSCHURLOCALUNKOWNLST, [37](#)
PASTIX_FORTRAN_GETSCHURLOCALUNKOWNNBR, [36](#)
PASTIX_FORTRAN_GETSCHUR, [35](#)
PASTIX_FORTRAN_SETBINDTAB, [28](#)
PASTIX_FORTRAN_SETCHURUNKNOWNLST, [35](#)
PASTIX_FORTRAN_SETSCHURARRAY, [37](#)
c_pastix, [38](#)
csc_dispatch, [31](#)
cscd_addlocal, [31](#)
cscd_checksymb, [30](#)
cscd_redispatch, [33](#)
cscd_symgraph, [30](#)
d_pastix, [38](#)
dpastix_fortran, [22](#)
dpastix, [22](#)
pastix_checkMatrix, [15](#), [29](#)
pastix_fortran, [20](#)
pastix_getLocalNodeLst, [27](#)
pastix_getLocalNodeNbr, [27](#)
pastix_getLocalUnknownLst, [28](#)
pastix_getLocalUnknownNbr, [27](#)
pastix_getSchurLocalNodeList, [35](#)
pastix_getSchurLocalNodeNbr, [35](#)
pastix_getSchurLocalUnkownList, [36](#)
pastix_getSchurLocalUnkownNbr, [36](#)
pastix_getSchur, [35](#)

- [pastix.setBind](#), [14](#), [28](#)
- [pastix.setSchurArray](#), [37](#)
- [pastix.setSchurUnknownList](#), [34](#)
- [pastix](#), [19](#)
- [s_pastix](#), [38](#)
- [z_pastix](#), [38](#)
- [SWIG](#), [39](#)
- [Threads](#), [9](#), [14](#)

Integer parameters

- [IPARM_ALLOCATED_TERMS](#), [17](#), [46](#)
- [IPARM_AMALGAMATION_LEVEL](#), [16](#)
- [IPARM_BINDTHRD](#), [14](#)
- [IPARM_CSCD_CORRECT](#), [14](#)
- [IPARM_DEFAULT_ORDERING](#), [15](#)
- [IPARM_DOF_NBR](#), [14](#)
- [IPARM_END_TASK](#), [13](#), [15](#), [18](#)
- [IPARM_ERROR_NUMBER](#), [46](#)
- [IPARM_ESP_NBTASKS](#), [46](#)
- [IPARM_FACTORISATION](#), [13](#)
- [IPARM_GMRES_IM](#), [18](#)
- [IPARM_INERTIA](#), [46](#)
- [IPARM_ITERMAX](#), [18](#)
- [IPARM_LEVEL_OF_FILL](#), [16](#)
- [IPARM_MATRIX_VERIFICATION](#), [14](#), [15](#)
- [IPARM_MAX_BLOCKSIZE](#), [16](#)
- [IPARM_METIS_PFACTOR](#), [16](#)
- [IPARM_MIN_BLOCKSIZE](#), [16](#)
- [IPARM_MODIFY_PARAMETER](#), [13](#), [15](#)
- [IPARM_NBITER](#), [18](#), [46](#)
- [IPARM_NNZEROS_BLOCK_LOCAL](#), [46](#)
- [IPARM_NNZERO](#), [16](#), [46](#)
- [IPARM_ORDERING_CMAX](#), [15](#)
- [IPARM_ORDERING_CMIN](#), [15](#)
- [IPARM_ORDERING_FRAT](#), [15](#)
- [IPARM_ORDERING_SWITCH_LEVEL](#), [15](#)
- [IPARM_ORDERING](#), [15](#)
- [IPARM_RHSD_CHECK](#), [41](#)
- [IPARM_RHS_MAKING](#), [17](#)
- [IPARM_SCHUR](#), [34](#)
- [IPARM_START_TASK](#), [13](#), [15](#), [18](#)
- [IPARM_STATIC_PIVOTING](#), [16](#), [17](#), [46](#)
- [IPARM_SYM](#), [14](#), [41](#)
- [IPARM_THREAD_NBR](#), [14](#)
- [IPARM_VERBOSE](#), [13](#), [44](#)

[MARCEL](#), [11](#)

[METIS](#), [15](#)

[Metis](#), [9](#), [11](#), [16](#)

[MPI](#), [9](#)

[Multi-arithmetic](#), [37](#), [41](#)

[PT-SCOTCH](#), [9](#), [11](#)

[python](#), [39](#)

[SCOTCH](#), [9](#), [11](#), [15](#)

Bibliography

- [AR12] Casadei Astrid and Pierre Ramet. Memory Optimization to Build a Schur Complement in an Hybrid Solver. Research Report RR-7971, INRIA, 2012.
- [BFL⁺12] G. Bosilca, M. Faverge, X. Lacoste, I. Yamazaki, and P. Ramet. Toward a supernodal sparse direct solver over DAG runtimes. In *Proceedings of PMAA'2012*, Londres, UK, June 2012.
- [BNP⁺06] B. Braconnier, B. Nkonga, M. Papin, P. Ramet, Ricchiuto M., J. Roman, and R. Abgrall. Efficient solution technique for low mach number compressible multiphase problems. In *Proceedings of PMAA'2006*, Rennes, France, September 2006.
- [CGR08] Y. Caniou, J.-S. Gay, and P. Ramet. Tunable parallel experiments in a gridrpc framework: application to linear solvers. In *VECPAR'08, 8th International Meeting High Performance Computing for Computational Science*, volume 5336 of *LNCS*, pages 430–436, Toulouse, France, June 2008. Springer Verlag.
- [Cha09] M. Chanaud. *Conception d'un solveur haute performance de systmes linaires creux couplant des mthodes multigrilles et directes pour la rsolution des quations de Maxwell 3D en rgime harmonique discrtes par lments finis*. PhD thesis, LaBRI, Universit Bordeaux I, Talence, Talence, France, December 2009.
- [CHR05] G. Caramel, P. Hénon, and P. Ramet. Etude de faisabilité pour la parallélisation d'un code de mécanique des fluides en version non structurée. Technical report, C.E.A. / C.E.S.T.A, 2005. Rapport Final.
- [CR07] G. Caramel and P. Ramet. Optimisation des performances des outils de calcul de neutronique des coeurs. Technical report, E.D.F. / SINETICS, 2007. Rapport Final.
- [CR12] A. Casadei and P. Ramet. Memory optimization to build a schur complement. In *SIAM Conference on Applied Linear Algebra*, Valence, Spain, June 2012.
- [CRR03] S. Christy, P. Ramet, and J. Roman. Dveloppement de la phase d'assemblage de la chane emilio pour un solveur parallle 2d. Technical report, C.E.A. / C.E.S.T.A, 2003. Rapport Final.
- [Fav09a] M. Faverge. Dynamic scheduling for sparse direct solver on numa and multicore architectures. In *Sparse Days*, Toulouse, France, June 2009.
- [Fav09b] M. Faverge. A numa aware scheduler for a parallel sparse direct solver. In *Journées Informatique Massivement Multiprocesseur et Multicoeur*, Rocquencourt, France, February 2009.
- [Fav09c] M. Faverge. *Ordonnancement hybride statique-dynamique en algre linéaire creuse pour de grands clusters de machines NUMA et multi-coeurs*. PhD thesis, LaBRI, Universit Bordeaux I, Talence, Talence, France, December 2009.

- [Fav09d] M. Faverge. Vers un solveur de systmes linaires creux adapt aux machines NUMA. In *ACTES RenPar'2009*, Toulouse, France, September 2009.
- [FLR08] M. Faverge, X. Lacoste, and P. Ramet. A numa aware scheduler for a parallel sparse direct solver. In *Proceedings of PMAA'2008*, Neuchatel, Swiss, June 2008.
- [FLR10] Mathieu Faverge, Xavier Lacoste, and Pierre Ramet. A NUMA Aware Scheduler for a Parallel Sparse Direct Solver. Technical report, 2010.
- [FR08] M. Faverge and P. Ramet. Dynamic scheduling for sparse direct solver on numa architectures. In *Proceedings of PARA'2008*, Trondheim, Norway, May 2008.
- [FR12] M. Faverge and P. Ramet. Fine grain scheduling for sparse solver on manycore architectures. In *15th SIAM Conference on Parallel Processing for Scientific Computing*, Savannah, USA, February 2012.
- [GHM⁺03] D. Goudin, P. Hénon, M. Mandallena, K. Mer, F. Pellegrini, P. Ramet, J. Roman, and J.-J. Pesqué. Outils numériques parallèles pour la résolution de très grands problèmes d'électromagnétisme. In *Séminaire sur l'Algorithmique Numérique Appliquée aux Problèmes Industriels*, Calais, France, May 2003.
- [GHP⁺99a] D. Goudin, P. Hénon, F. Pellegrini, P. Ramet, J. Roman, and J.-J. Pesqué. Algèbre linéaire creuse parallèle pour les méthodes directes : Application à la parallélisation d'un code de mécanique des structures. In *Journées sur l'Algèbre Linéaire Creuse et ses Applications Industrielles*, Rennes, France, 1999.
- [GHP⁺99b] D. Goudin, P. Hénon, F. Pellegrini, P. Ramet, J. Roman, and J.-J. Pesqué. Parallel sparse linear algebra and application to structural mechanics. In *SPWorld'99*, Montpellier, France, 1999.
- [GHP⁺00a] D. Goudin, P. Hénon, F. Pellegrini, P. Ramet, and J. Roman. Résolution parallèle de grands systèmes linéaires creux. In *Proceedings of JSFT'2000*, Monastir, Tunisia, October 2000.
- [GHP⁺00b] D. Goudin, P. Hénon, F. Pellegrini, P. Ramet, J. Roman, and J.-J. Pesqué. Algèbre Linéaire Creuse Hautes Performances : Application à la Mécanique des Structures. In *iHPerf'2000*, Aussois, France, December 2000.
- [GHP⁺00c] D. Goudin, P. Hénon, F. Pellegrini, P. Ramet, J. Roman, and J.-J. Pesqué. Description of the emilio software processing chain and application to structural mechanics. In *Proceedings of PMAA'2K*, Neuchatel, Swiss, August 2000.
- [GHP⁺00d] D. Goudin, P. Hénon, F. Pellegrini, P. Ramet, J. Roman, and J.-J. Pesqué. Parallel sparse linear algebra and application to structural mechanics. *Numerical Algorithms*, 24:371–391, 2000.
- [GHP⁺00e] D. Goudin, P. Hénon, F. Pellegrini, P. Ramet, J. Roman, and J.-J. Pesqué. Parallel sparse linear algebra and application to structural mechanics. In *European ACTC Workshop*, Paris, France, May 2000.
- [GHP⁺01] D. Goudin, P. Hénon, F. Pellegrini, P. Ramet, and J. Roman. Mise en oeuvre d'une bibliothèque d'outils pour la résolution par méthode directe de grands systèmes linéaires creux symétriques définis positifs sur machine parallèle. Technical report, C.E.A. / C.E.S.T.A, 2001. Manuel utilisateur de la chaîne EMILIO.
- [Gou98] D. Goudin. Mise en oeuvre d'une bibliothèque d'outils pour la résolution par méthode directe de grands systèmes linéaires creux symétriques définis positifs sur machine parallèle. Technical report, C.E.A. / C.E.S.T.A, 1998. Premier Rapport Semestriel.
- [Gou99a] D. Goudin. Mise en oeuvre d'une bibliothèque d'outils pour la résolution par méthode directe de grands systèmes linéaires creux symétriques définis positifs sur machine

- parallèle. Technical report, C.E.A. / C.E.S.T.A, 1999. Rapport Final de la Première Partie.
- [Gou99b] D. Goudin. Mise en oeuvre d'une bibliothèque d'outils pour la résolution par méthode directe de grands systèmes linéaires creux symétriques définis positifs sur machine parallèle. Technical report, C.E.A. / C.E.S.T.A, 1999. Rapport Semestriel de la Deuxième Partie.
- [Gou00a] D. Goudin. Assemblage parallèle d'une matrice et/ou d'un second membre: Application à la parallélisation d'un code de mécanique des structures. In *ACTES Ren-Par'2000*, Besancon, France, 2000.
- [Gou00b] D. Goudin. *Mise en œuvre d'une Bibliothèque d'Outils pour la Résolution Parallèle Hautes Performances par Méthode Directe de Grands Systèmes Linéaires Creux et application à un Code de Mécanique des Structures*. PhD thesis, LaBRI, Universit Bordeaux I, Talence, France, November 2000.
- [Gou00c] D. Goudin. Mise en oeuvre d'une bibliothèque d'outils pour la résolution par méthode directe de grands systèmes linéaires creux symétriques définis positifs sur machine parallèle. Technical report, C.E.A. / C.E.S.T.A, 2000. Rapport Final de la Deuxième Partie.
- [GR00] D. Goudin and J. Roman. A scalable parallel assembly for irregular meshes based on a block distribution for a parallel block direct solver. In *Proceedings of PARA'2000*, volume 1947 of *LNCS*, Bergen, Norway, 2000. Springer Verlag.
- [GRR04] A. Gouremen, P. Ramet, and J. Roman. Dveloppement de la phase d'assemblage de la chane emilio (distribution du maillage et multi-threading). Technical report, C.E.A. / C.E.S.T.A, 2004. Rapport Final.
- [H01] P. Hénon. *Distribution des Données et Régulation Statique des Calculs et des Communications pour la Résolution de Grands Systèmes Linéaires Creux par Méthode Directe*. PhD thesis, LaBRI, Universit Bordeaux I, Talence, Talence, France, November 2001.
- [HHR05] F. Huard, P. Hénon, and P. Ramet. Intégration dans odyssee de la chaine logicielle emilio. Technical report, C.E.A. / C.E.S.T.A, 2005. Rapport Final.
- [HLRR03] P. Hénon, D. Lecas, P. Ramet, and J. Roman. Amélioration et extension du solveur direct parallèle pour grandes matrices creuses du cesta. Technical report, C.E.A. / C.E.S.T.A, 2003. Rapport Final.
- [HNRR04] P. Hénon, B. Nkonga, P. Ramet, and J. Roman. Using of the high performance sparse solver pastix for the complex multiscale 3d simulations performed by the fluidbox fluid mechanics software. In *Proceedings of PMAA'2004*, Marseille, France, October 2004.
- [HPR⁺04a] P. Hénon, F. Pellegrini, P. Ramet, J. Roman, and Y. Saad. Applying parallel direct solver skills to build robust and highly performant preconditioners. In *Proceedings of PARA'2004*, volume 3732 of *LNCS*, pages 601–619, Copenhagen, Denmark, June 2004. Springer Verlag.
- [HPR⁺04b] P. Hénon, F. Pellegrini, P. Ramet, J. Roman, and Y. Saad. High performance complete and incomplete factorizations for very large sparse systems by using Scotch and PaStiX softwares. In *Eleventh SIAM Conference on Parallel Processing for Scientific Computing*, San Francisco, USA, February 2004.
- [HPRR03a] P. Hénon, F. Pellegrini, P. Ramet, and J. Roman. An efficient hybrid mpi/thread implementation on a network of smp nodes for the parallel sparse direct solver pastix: ordering / scheduling / memory managment / out-of-core issues, and application to preconditioning. In *Sparse Days*, Saint Girons, France, June 2003.

- [HPRR03b] P. Hénon, F. Pellegrini, P. Ramet, and J. Roman. Towards high performance hybrid direct-iterative solvers for large sparse systems. In *International SIAM Conference On Preconditioning Techniques For Large Sparse Matrix Problems In Scientific And Industrial Applications*, Napa Valley, USA, October 2003.
- [HPRR04] P. Hénon, F. Pellegrini, P. Ramet, and J. Roman. Etude sur l'applicabilité de méthodes itératives nouvelles aux problèmes du cesta. Technical report, C.E.A. / C.E.S.T.A., 2004. Rapport Final.
- [HPRR05] P. Hénon, F. Pellegrini, P. Ramet, and J. Roman. Blocking issues for an efficient parallel block ilu preconditioner. In *International SIAM Conference On Preconditioning Techniques For Large Sparse Matrix Problems In Scientific And Industrial Applications*, Atlanta, USA, May 2005.
- [HR01] P. Hénon and P. Ramet. PaStiX: Un solveur parallèle direct pour des matrices creuses symétriques définies positives basé sur un ordonnancement statique performant et sur une gestion mémoire efficace. In *ACTES RenPar'2001*, Paris, France, April 2001.
- [HR02] P. Hénon and P. Ramet. Optimisation de l'occupation mémoire pour un solveur parallèle creux direct hautes performances de type supernodal. In *ACTES RenPar'2002*, Hamamet, Tunisia, April 2002.
- [HR10a] P. Hénon and P. Ramet. Scalable direct and iterative solvers. In *SuperComputing'2010*, New Orleans, USA, November 2010.
- [HR10b] P. Hénon and P. Ramet. Scalable direct and iterative solvers, June 2010. Workshop INRIA-UUIC, Bordeaux, France.
- [HRR99] P. Hénon, P. Ramet, and J. Roman. A Mapping and Scheduling Algorithm for Parallel Sparse Fan-In Numerical Factorization. In *Proceedings of Euro-Par'99*, volume 1685 of *LNCS*, pages 1059–1067, Toulouse, France, September 1999. Springer Verlag.
- [HRR00a] P. Hénon, P. Ramet, and J. Roman. Pastix: A high-performance parallel direct solver for sparse symmetric definite systems. In *Proceedings of PMAA '2K*, Neuchatel, Swiss, August 2000.
- [HRR00b] P. Hénon, P. Ramet, and J. Roman. PaStiX: A Parallel Sparse Direct Solver Based on a Static Scheduling for Mixed 1D/2D Block Distributions. In *Proceedings of Irregular'2000 workshop of IPDPS*, volume 1800 of *LNCS*, pages 519–525, Cancun, Mexico, May 2000. Springer Verlag.
- [HRR01] P. Hénon, P. Ramet, and J. Roman. PaStiX: A Parallel Direct Solver for Sparse SPD Matrices based on Efficient Static Scheduling and Memory Management. In *Tenth SIAM Conference on Parallel Processing for Scientific Computing*, Portsmouth, USA, March 2001.
- [HRR02a] P. Hénon, P. Ramet, and J. Roman. A Parallel Direct Solver for Very Large Sparse SPD Systems. In *SuperComputing'2002*, Baltimore, USA, November 2002.
- [HRR02b] P. Hénon, P. Ramet, and J. Roman. Parallel factorization of very large sparse SPD systems on a network of SMP nodes. In *Proceedings of PMAA '2002*, Neuchatel, Swiss, November 2002.
- [HRR02c] P. Hénon, P. Ramet, and J. Roman. PaStiX: A High-Performance Parallel Direct Solver for Sparse Symmetric Definite Systems. *Parallel Computing*, 28(2):301–321, January 2002.
- [HRR03a] P. Hénon, P. Ramet, and J. Roman. A Parallel Direct Solver for Very Large Sparse SPD Systems. In *IPDPS'2003*, Nice, France, April 2003.

- [HRR03b] P. Hénon, P. Ramet, and J. Roman. Efficient algorithms for direct resolution of large sparse system on clusters of SMP nodes. In *SIAM Conference on Applied Linear Algebra*, Williamsburg, USA, July 2003.
- [HRR04a] P. Hénon, P. Ramet, and J. Roman. A blockwise algorithm for parallel incomplete cholesky factorization. In *Proceedings of PMAA'2004*, Marseille, France, October 2004.
- [HRR04b] P. Hénon, P. Ramet, and J. Roman. Parallel Complete and Incomplete Blockwise Factorisations for Very Large Sparse Systems. In *SuperComputing'2004*, Pittsburgh, USA, November 2004.
- [HRR05a] P. Hénon, P. Ramet, and J. Roman. Evaluation des performances de la version smp du solveur pastix de la chaine logicielle emilio dans l'environnement du code odyssee. Technical report, C.E.A. / C.E.S.T.A, 2005. Rapport Final.
- [HRR05b] P. Hénon, P. Ramet, and J. Roman. On using an hybrid mpi-thread programming for the implementation of a parallel sparse direct solver on a network of smp nodes. In *Proceedings of Sixth International Conference on Parallel Processing and Applied Mathematics, Workshop HPC Linear Algebra*, volume 3911 of *LNCS*, pages 1050–1057, Poznan, Poland, September 2005. Springer Verlag.
- [HRR06a] P. Hénon, P. Ramet, and J. Roman. On finding approximate supernodes for an efficient ilu(k) factorization. In *Proceedings of PMAA'2006*, Rennes, France, September 2006.
- [HRR06b] P. Hénon, P. Ramet, and J. Roman. Partitioning and blocking issues for a parallel incomplete factorization. In *Proceedings of PARA'2006*, volume 4699 of *LNCS*, pages 929–937, Umea, Sweden, June 2006. Springer Verlag.
- [HRR07] P. Hénon, P. Ramet, and J. Roman. A supernode amalgamation algorithm for an efficient block incomplete factorization. In *Proceedings of PPAM'2007, CTPSM07 Workshop*, Gdansk, Poland, September 2007.
- [HRR08a] P. Hénon, P. Ramet, and J. Roman. On finding approximate supernodes for an efficient ilu(k) factorization. *Parallel Computing*, 34:345–362, 2008.
- [HRR08b] P. Hénon, P. Ramet, and J. Roman. A supernode amalgamation algorithm for an efficient block incomplete factorization, July 2008. Mini-Workshop on parallel iterative solvers and domain decomposition techniques, Minneapolis, USA.
- [KST⁺08] N. Kushida, Y. Suzuki, N. Teshima, N. Nakajima, Y. Caniou, M. Dayde, and P. Ramet. Toward an international sparse linear algebra expert system by inter-connecting the itbl computational grid with the grid-tlse platform. In *VECPAR'08, 8th International Meeting High Performance Computing for Computational Science*, volume 5336 of *LNCS*, pages 424–429, Toulouse, France, June 2008. Springer Verlag.
- [LFR12a] X. Lacoste, M. Faverge, and P. Ramet. Scheduling for sparse solver on manycore architectures. In *Workshop INRIA-CNPq, HOSCAR meeting*, Petropolis, Brazil, September 2012.
- [LFR12b] X. Lacoste, M. Faverge, and P. Ramet. Sparse direct solvers with accelerators over DAG runtimes. In *Workshop INRIA-CNPq, HOSCAR meeting*, Sophia-Antipolis, France, July 2012.
- [LR12] X. Lacoste and P. Ramet. Sparse direct solver on top of large-scale multicore systems with gpu accelerators. In *SIAM Conference on Applied Linear Algebra*, Valence, Spain, June 2012.

- [LRF⁺12] Xavier Lacoste, Pierre Ramet, Mathieu Faverge, Yamazaki Ichitaro, and Jack Dongarra. Sparse direct solvers with accelerators over DAG runtimes. Research Report RR-7972, INRIA, 2012.
- [Ram00] P. Ramet. *Optimisation de la Communication et de la Distribution des Données pour des Solveurs Parallèles Directs en Algèbre Linéaire Dense et Creuse*. PhD thesis, LaBRI, Université Bordeaux I, Talence, France, January 2000.
- [Ram07] P. Ramet. High performances methods for solving large sparse linear systems - direct and incomplete factorization. In *Second NExt Grid Systems and Techniques, REDIMSPS Workshop*, Tokyo, Japan, May 2007.
- [Ram09] P. Ramet. Dynamic scheduling for sparse direct solver on numa and multicore architectures. In *ComplexHPC meeting*, Lisbon, Portugal, October 2009.
- [RR01] P. Ramet and J. Roman. Analyse et étude de faisabilité de la résolution par méthode directe sur machine parallèle de grands systèmes linéaires symétriques définis positifs pour des problèmes d'électromagnétisme avec couplage éléments finis – équations intégrales. Technical report, C.E.A. / C.E.S.T.A, 2001. Rapport Final.
- [Saa96] Y. Saad. *Iterative Methods For Sparse Linear Systems*. Ed. PWS publishing Company, 1996.
- [SKT⁺10] Y. Suzuki, N. Kushida, T. Tatekawa, N. Teshima, Y. Caniou, R. Guivarch, M. Dayde, and P. Ramet. Development of an International Matrix-Solver Prediction System on a French-Japanese International Grid Computing Environment. In *Joint International Conference on Supercomputing in Nuclear Applications and Monte Carlo 2010 (SNA + MC2010)*, Tokyo, Japan, October 2010.